



Introduction to Spark Streaming

Surendra Panpaliya

Agenda



What is Spark streaming?



High level streaming Architecture



Spark streaming Sources



Spark streaming's place in spark



Spark structured streaming place in spark

Agenda



Dstream, Transformations



Discretized Steams (DStreams)



Receivers



Batching



Dstream processing

Agenda

Spark structured streaming

How does it work?

Steps for structured streaming

Supported sources and sinks

What is Spark Streaming?



An extension of the
core Spark API



enables scalable,



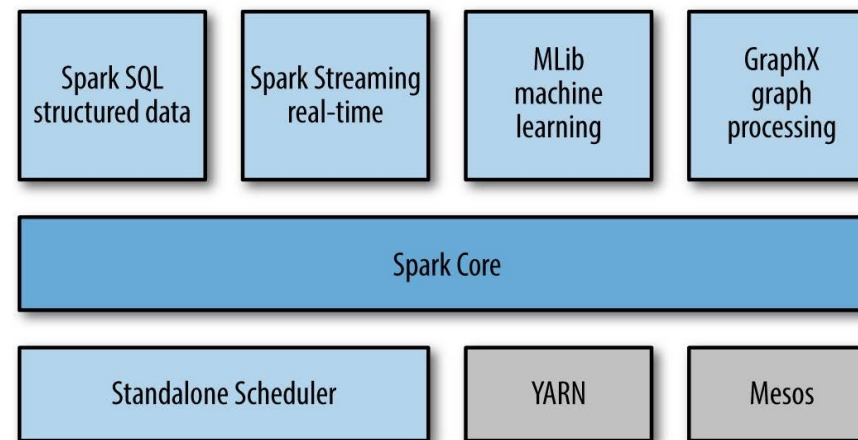
high-throughput,



fault-tolerant



stream processing of
live data streams.





What is Spark Streaming?

- Data can be ingested
- from many sources like
- Kafka, Kinesis,
- TCP sockets



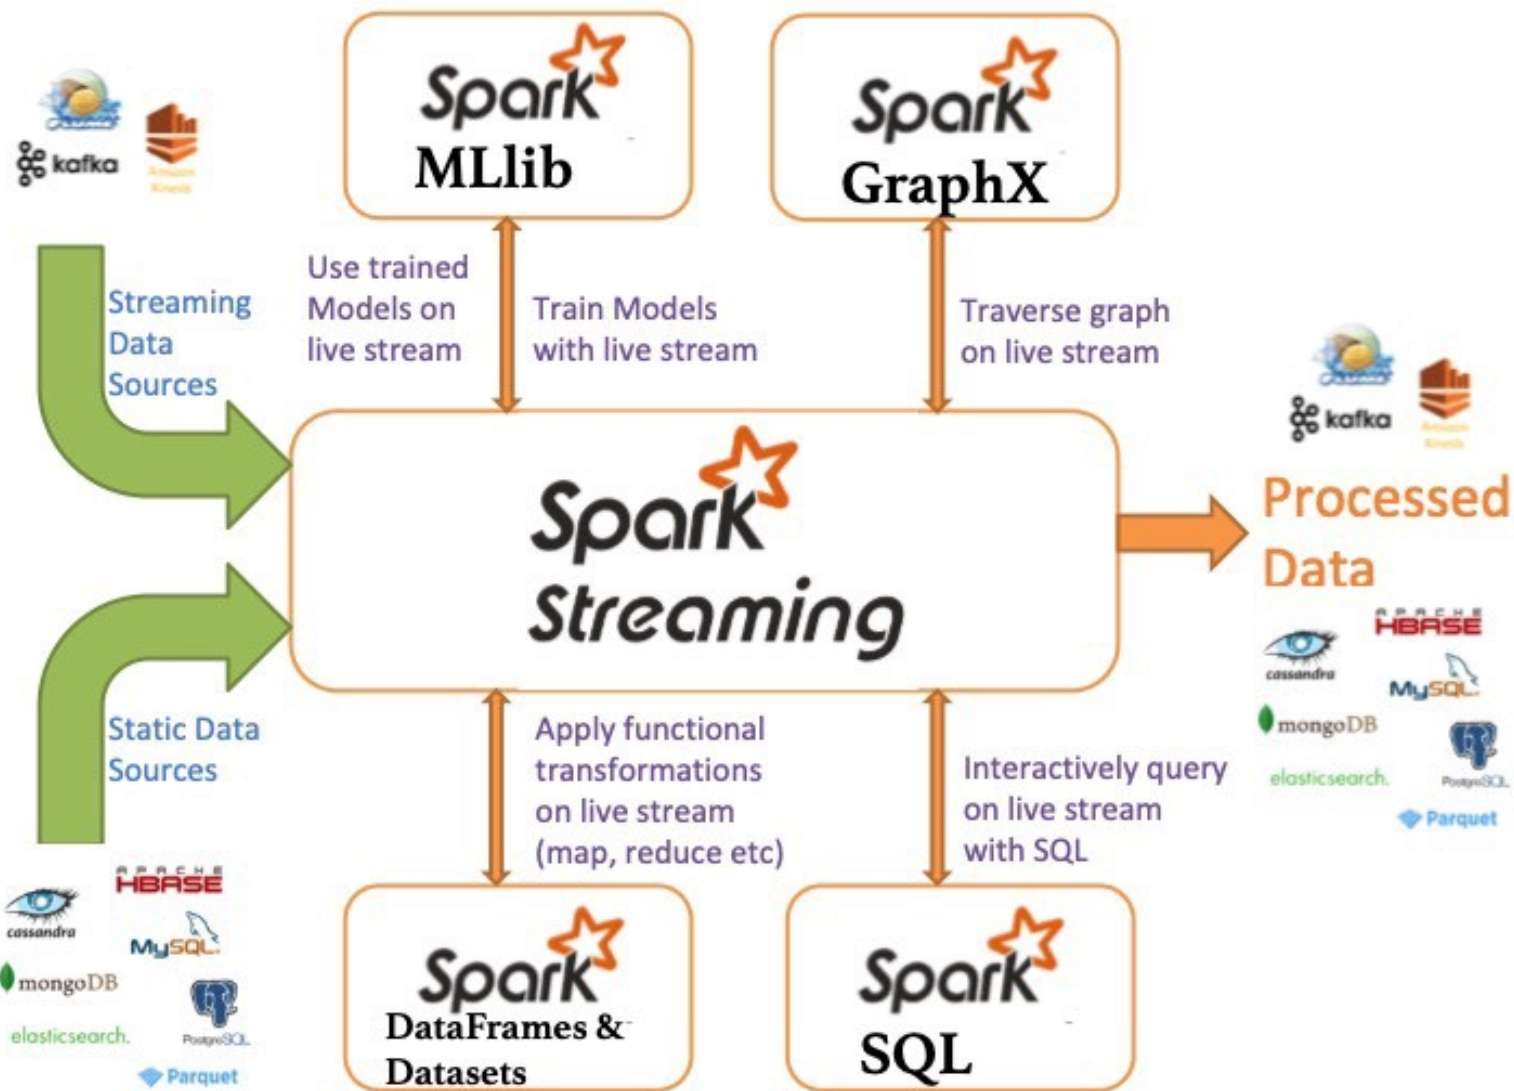
What is Spark Streaming?

- Can be processed
- using complex algorithms
- with high-level functions like
- map, reduce, join and window.



What is Spark Streaming?

- Finally, processed data
- can be pushed out to
- filesystems, databases, and
- live dashboards.



What is Spark Streaming?



Can apply Spark's



machine learning and



graph processing algorithms



on data streams.



How it works?

- Spark Streaming
- receives live input data streams and
- divides the data into batches




How it works?

- processed by the Spark engine
- to generate the
- final stream of results in batches.

High level streaming Architecture

Spark Streaming

discretized stream processing



records processed in batches with short tasks
each batch is a RDD (partitioned dataset)

High level streaming Architecture

At a high level, modern distributed stream processing pipelines execute as follows:

Receive streaming data from data sources


(e.g. live logs, system telemetry data, IoT device data, etc.)

into some data ingestion system like

Apache Kafka, Amazon Kinesis, etc.

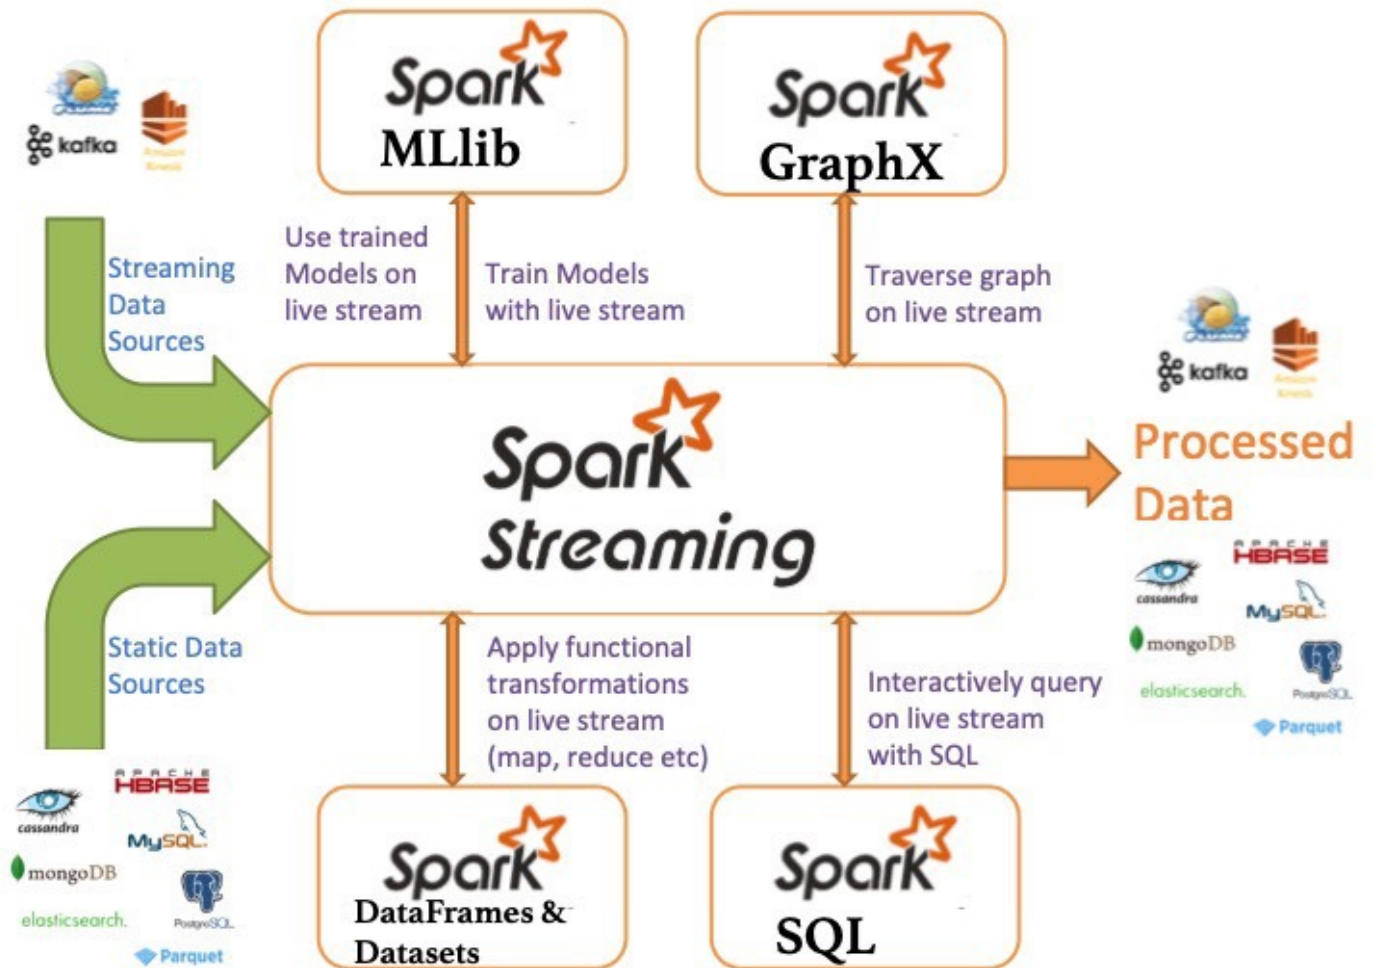
High level streaming Architecture

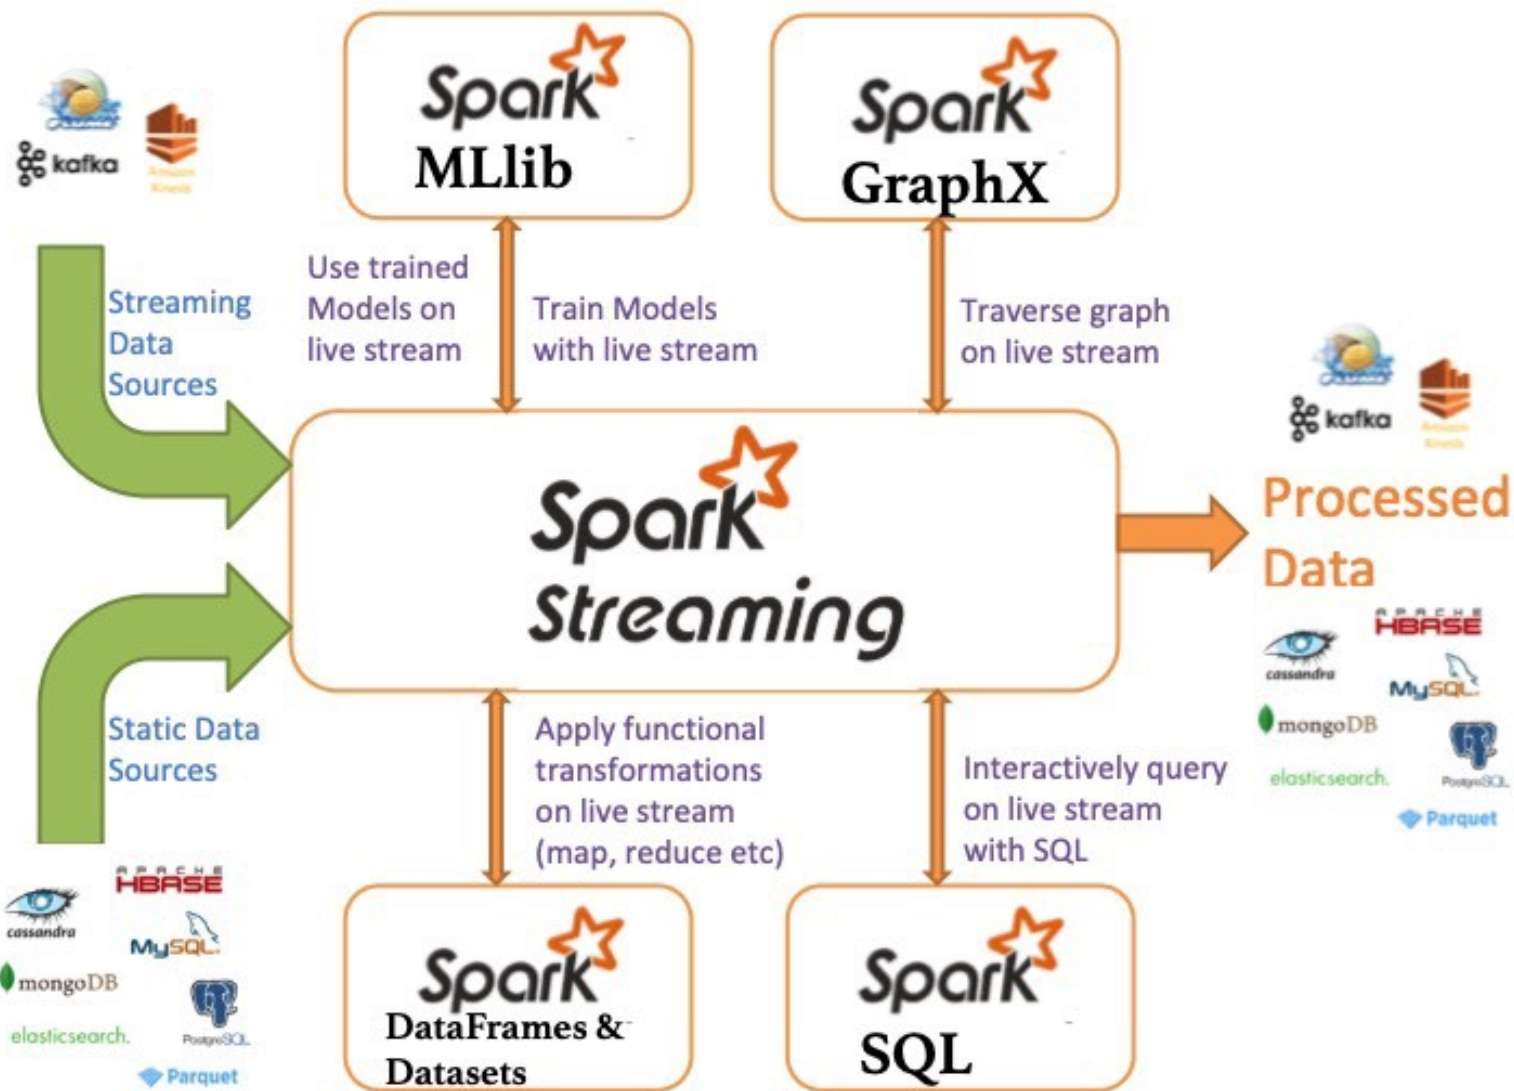
- Stream Process Engine
- **Process** the data in
- parallel on a cluster.



High level streaming Architecture

- **Output** the results
- out to downstream systems like
- HBase, Cassandra, Kafka, etc.





Spark streaming Sources



File Streams:



Reading data from files



on any file system



compatible with the HDFS API



HDFS, S3, NFS



```
DStream = streamingContext.fileStream<...>(directory);
```

Spark streaming Sources

Streams based on Custom Receivers:

DStreams can be created

with data streams received

through custom receivers,

extending the `Receiver<T>` class...

```
streamingContext.queueStream(queueOfRDDs)
```

Spark streaming Sources

Queue of RDDs as a Stream:

For testing a Spark Streaming application

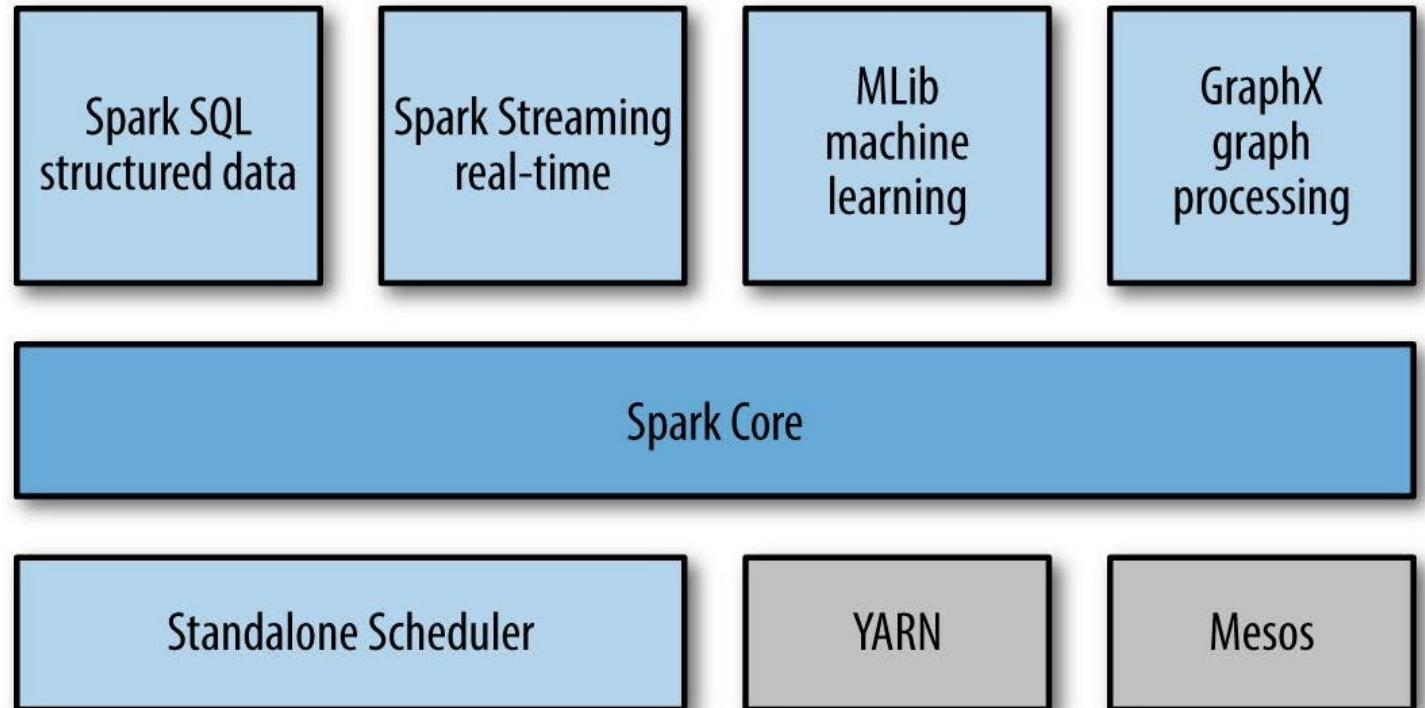
with test data

DStream based on a queue of RDDs

```
streamingContext.queueStream(queueOfRDDs)
```

Spark streaming's place in spark

- **Spark Streaming:**
- A component that
- enables processing of
- live streams of data



Spark structured streaming place in spark



The Spark SQL engine



will take care of running it incrementally

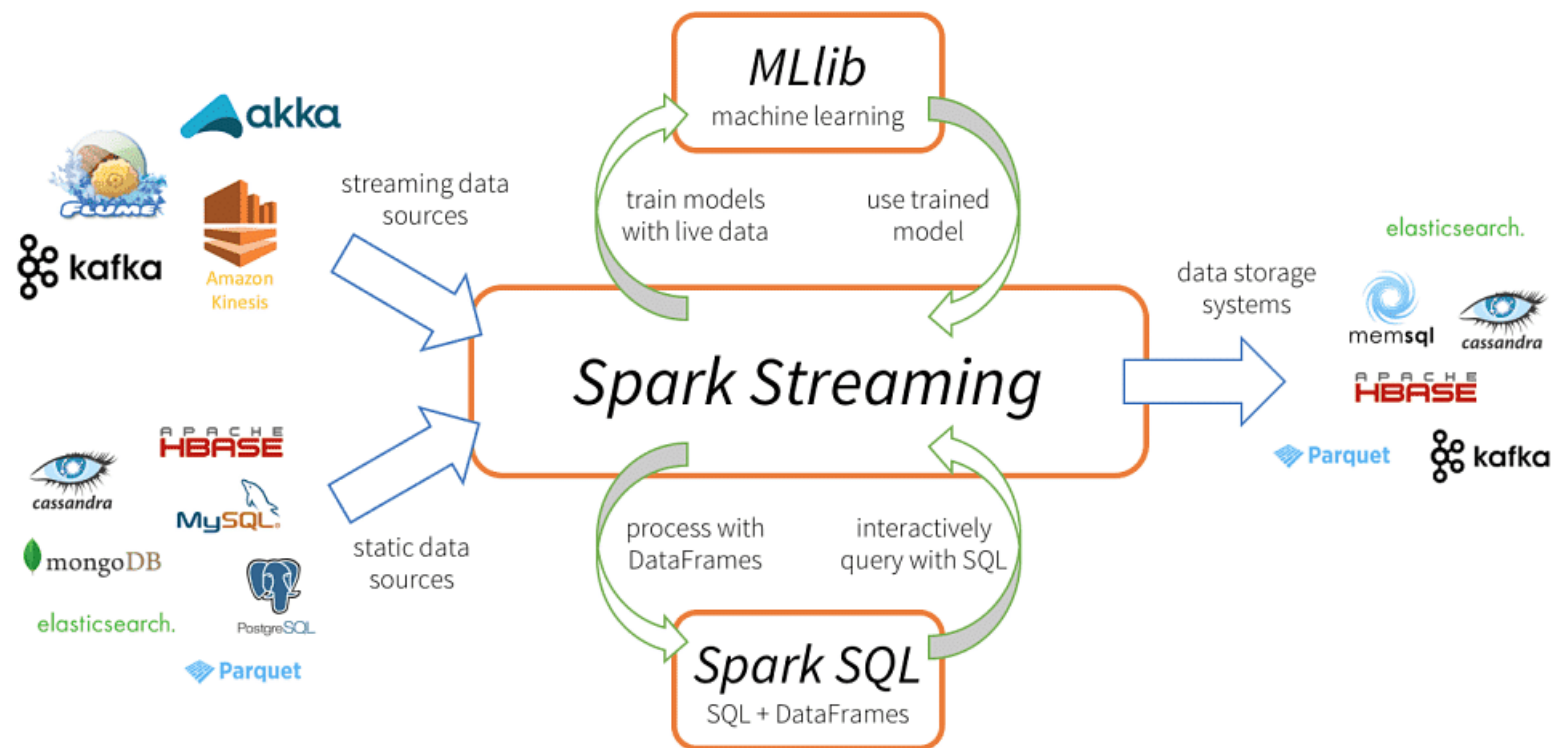


Continuously and updating the final result




as streaming data continues to arrive.

Spark structured streaming




Structured Streaming is a scalable and fault-tolerant stream processing engine built on the Spark SQL engine.



Streams as tables


- Spark Structured Streaming represents
- a stream of data as a table





Streams as tables

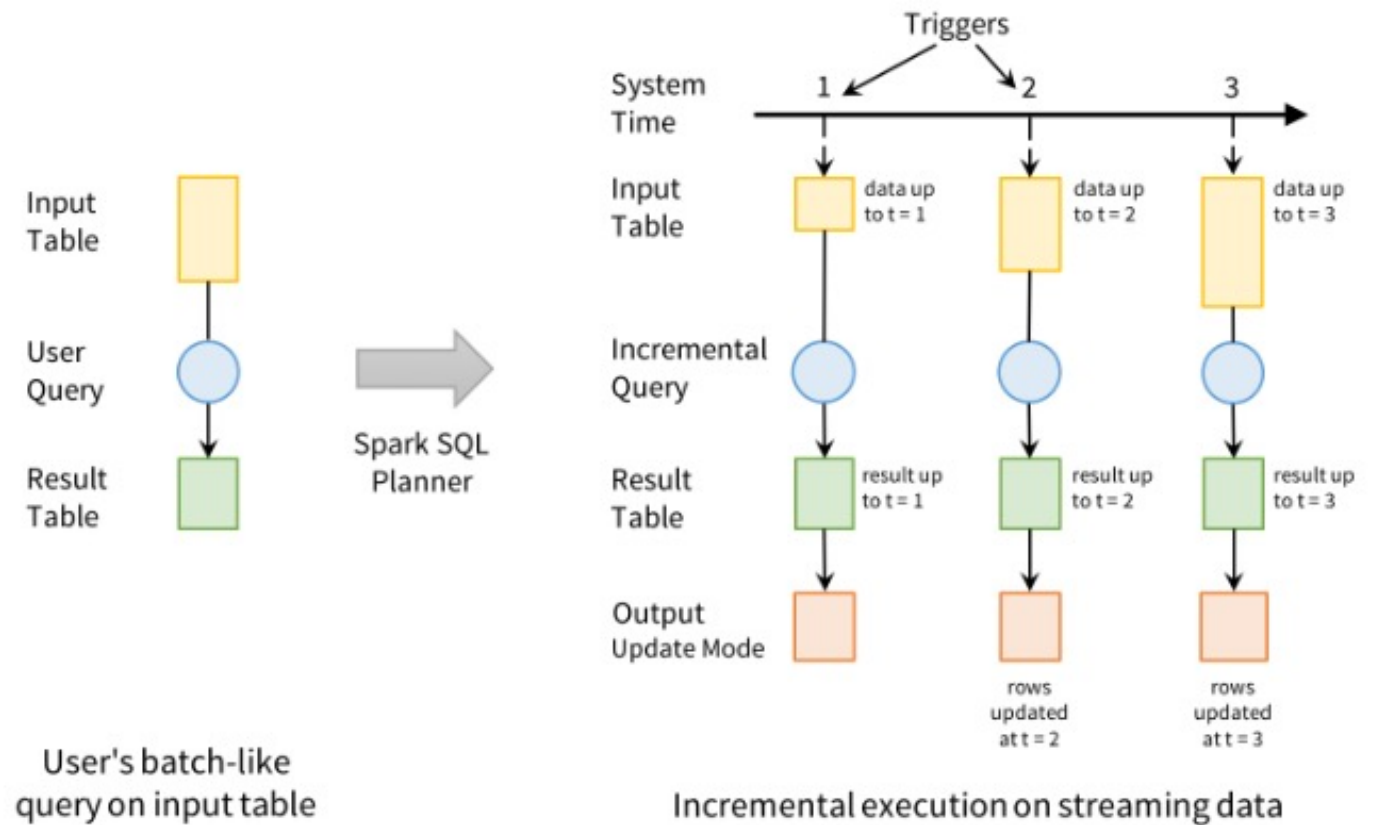
- Unbounded in depth,
- the table continues to grow as new data arrives.



Streams as tables


- This *input table* is continuously processed
- by a long-running query
- the results sent to an *output table*

Spark structured streaming




Structured Streaming Processing Model

Users express queries using a batch API; Spark incrementalizes them to run on streams




Discretized Streams (DStreams)

- Basic abstraction provided by Spark Streaming.




Discretized Streams (DStreams)

- Continuous stream of data,
- input data stream received from source




Discretized Streams (DStreams)

- The processed data stream
- generated by transforming
- the input stream.



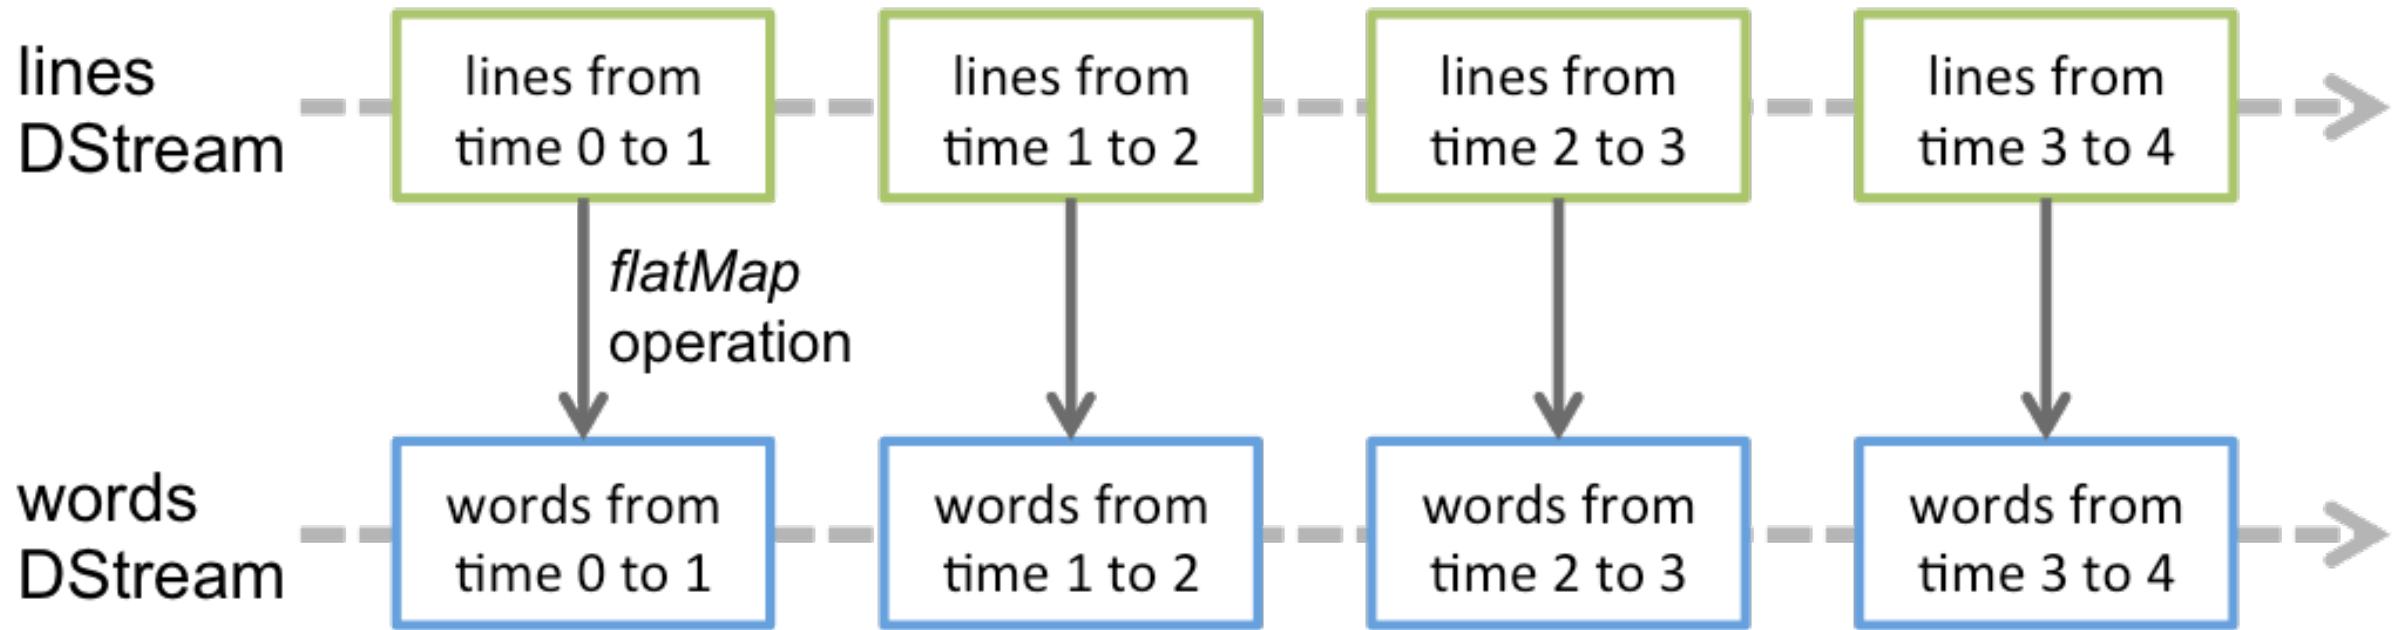
Discretized Streams (DStreams)

- Internally, a DStream is represented
- by a continuous series of RDDs




Discretized Streams (DStreams)

- Internally, a DStream is represented
- by a continuous series of RDDs
- Abstraction of an immutable,
- distributed dataset



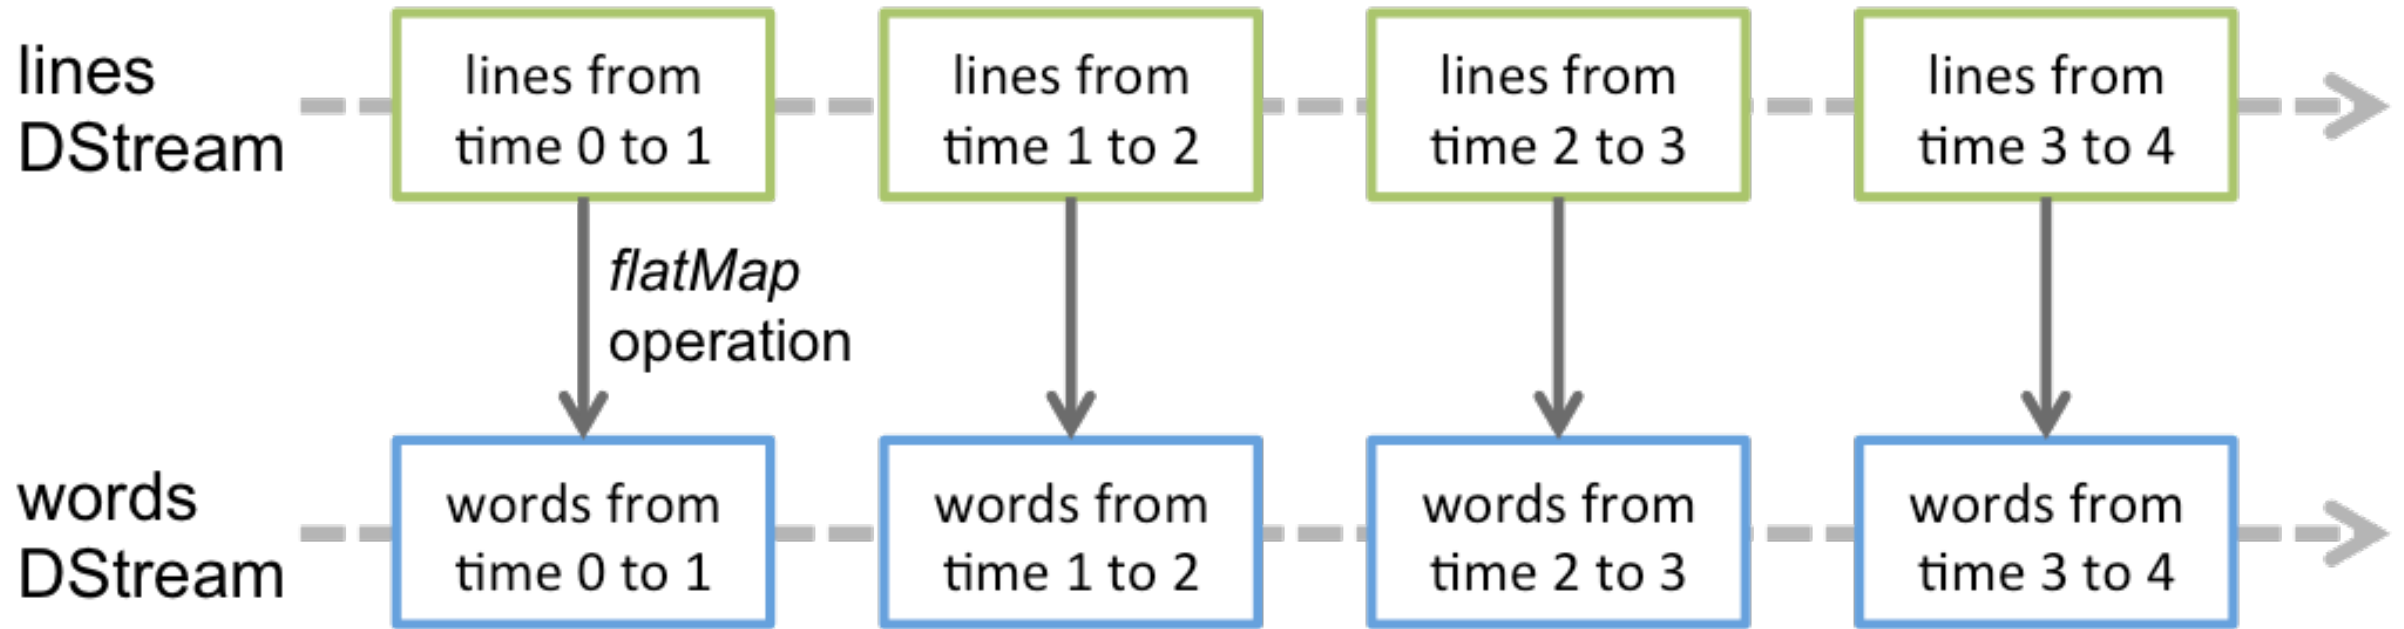
Discretized Streams (DStreams)

- Each RDD in a DStream
- contains data from
- a certain interval.



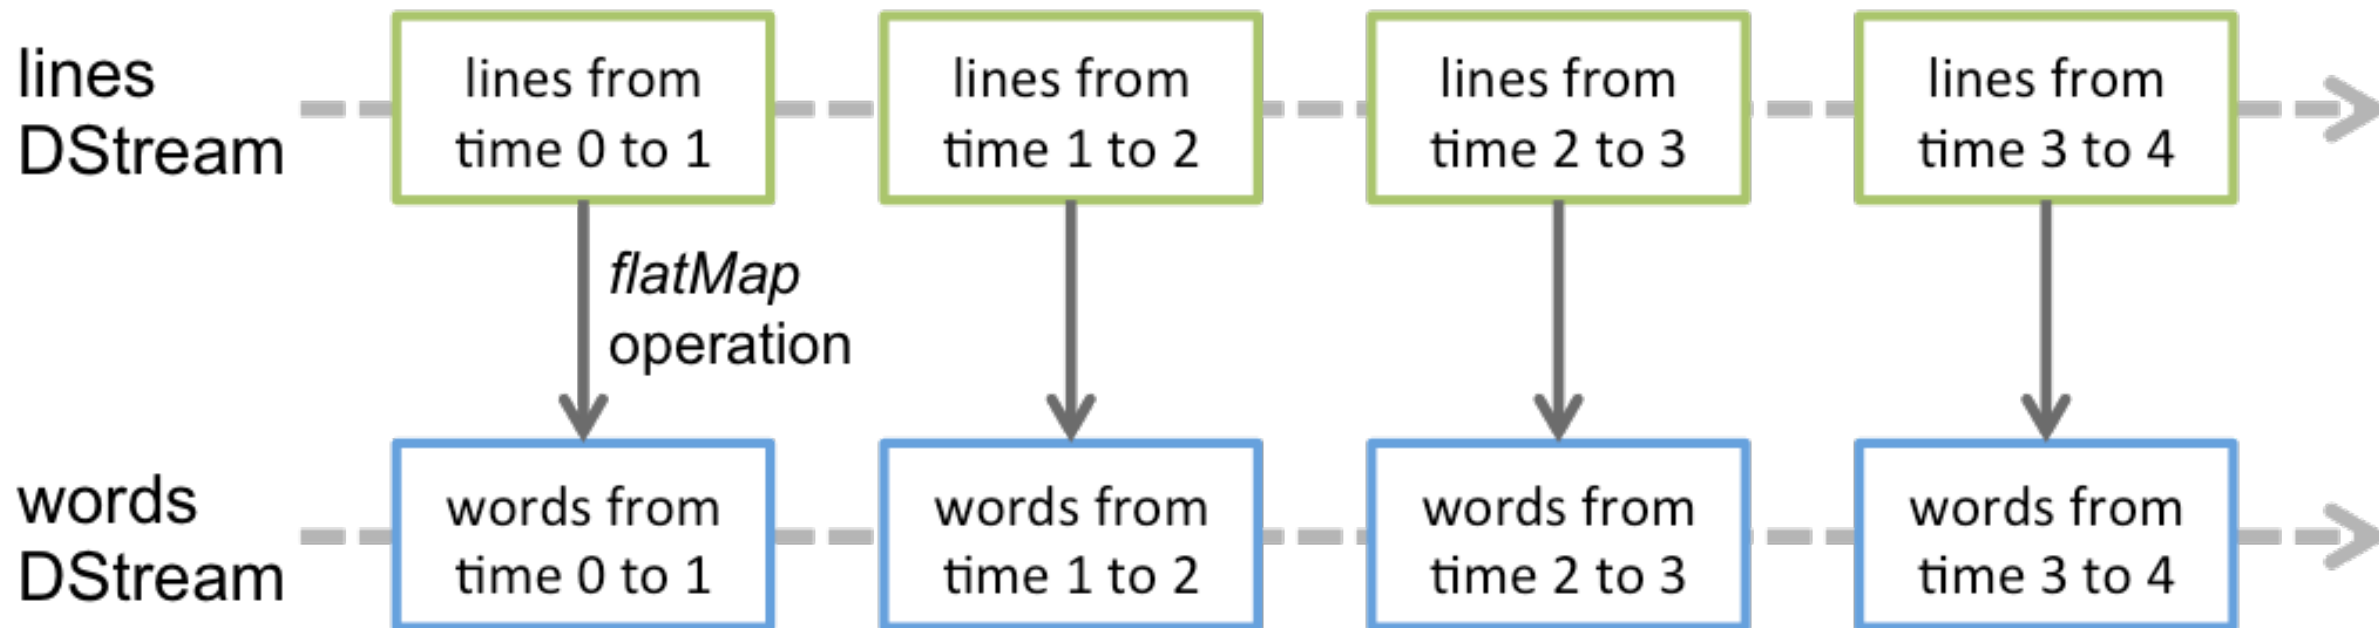
Discretized Streams (DStreams)

- Any operation applied on a
- DStream translates
- to operations on the underlying RDDs



Discretized Streams (DStreams)

- Converting a stream of lines to words,
- the `flatMap` operation is applied on each RDD
- in the `lines DStream` to generate
- the RDDs of the `words DStream`.



Discretized Streams (DStreams)

- Underlying RDD transformations
- are computed by the Spark engine
- The DStream operations hide most of these details
- Provide the developer with a higher-level API for convenience.

Input DStreams and Receivers

Input DStreams are
DStreams

representing the stream
of input data

received from streaming
sources.

Input DStreams and Receivers



lines was an input DStream as



it represented the stream of data



received from the netcat server.

Input DStreams and Receivers

Every input DStream

is associated with

a **Receiver** (Scala) object

which receives

the data from a source

stores it in Spark's memory

for processing.

A sample program

Spark structured streaming



Provides fast,



scalable,



fault-tolerant,



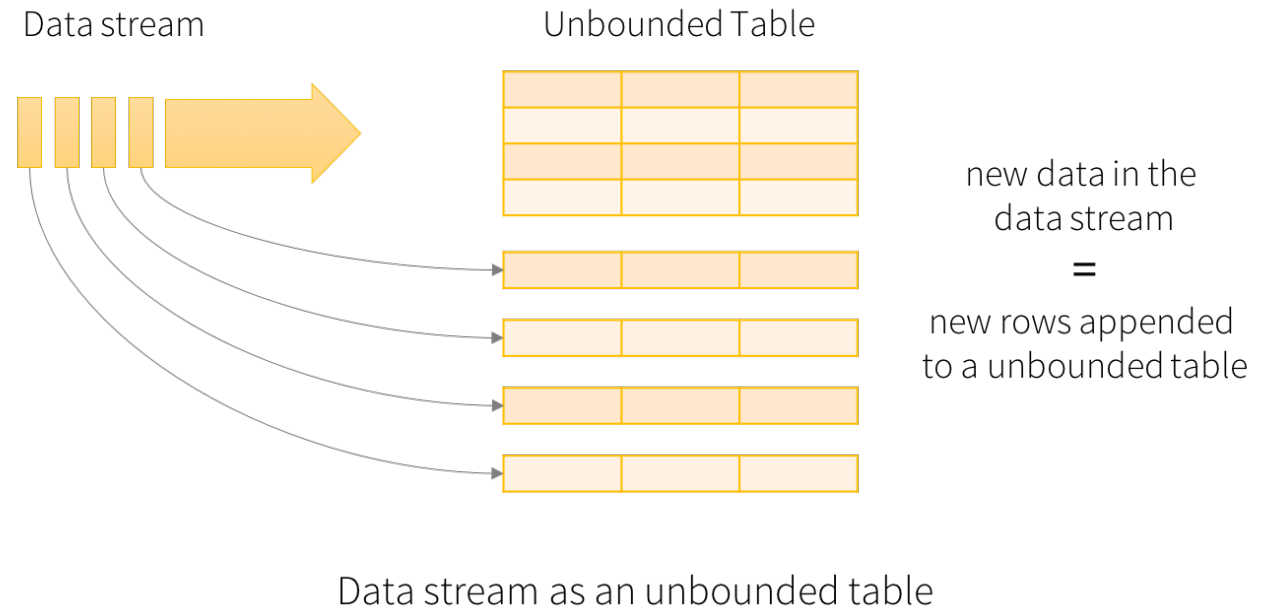
end-to-end exactly-once



stream processing

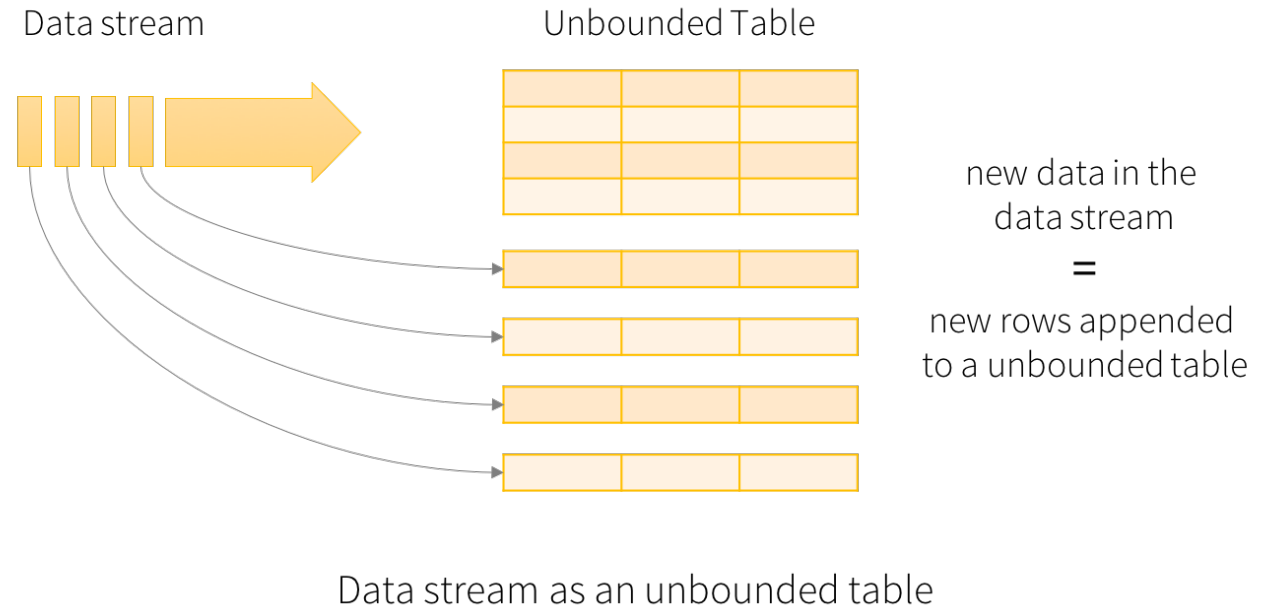
How does it work?

- Structured Streaming queries are
- processed using
- a *micro-batch processing engine*



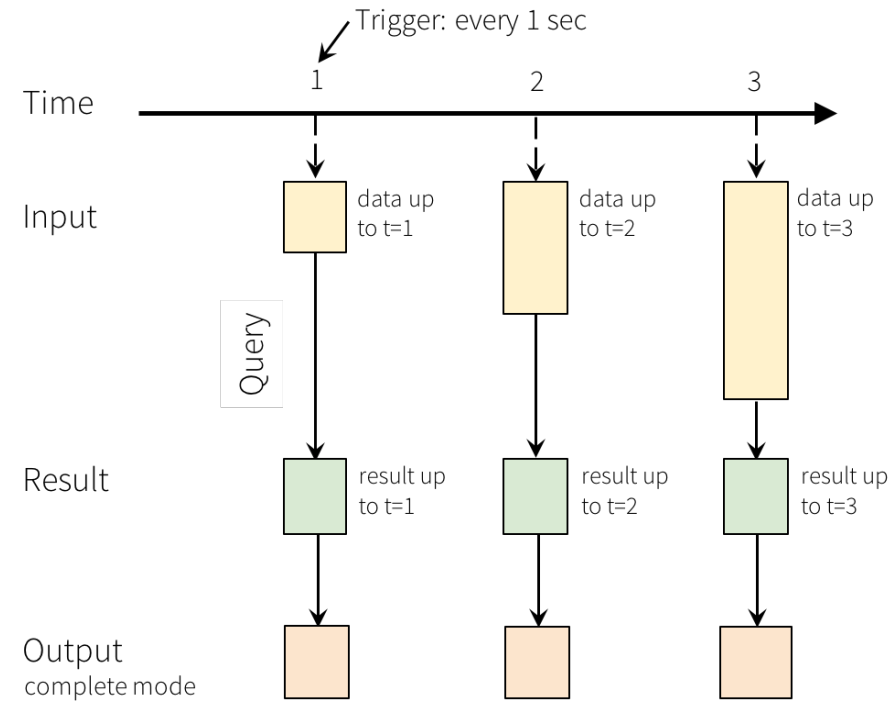
How does it work?

- which processes data streams
- as a series of small batch jobs



How does it work?


- achieving end-to-end latencies
- as low as 100 milliseconds and
- exactly-once fault-tolerance guarantees.



Programming Model for Structured Streaming

How does it work?

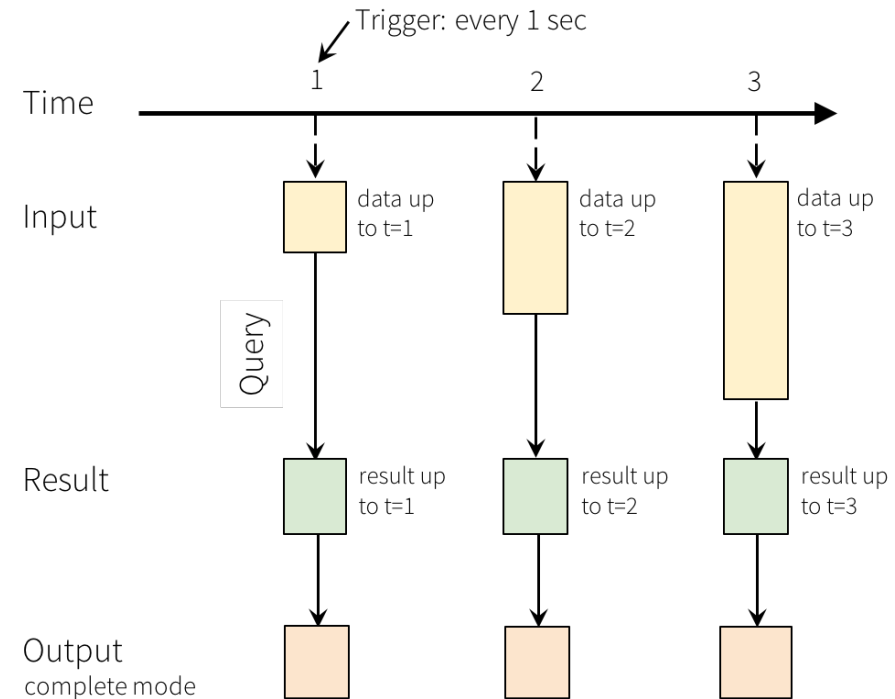
- In Spark 2.3 introduced
- a new low-latency processing mode
- called **Continuous Processing**



Programming Model for Structured Streaming

How does it work?


- which can achieve end-to-end latencies
- as low as 1 millisecond
- with at-least-once guarantees.



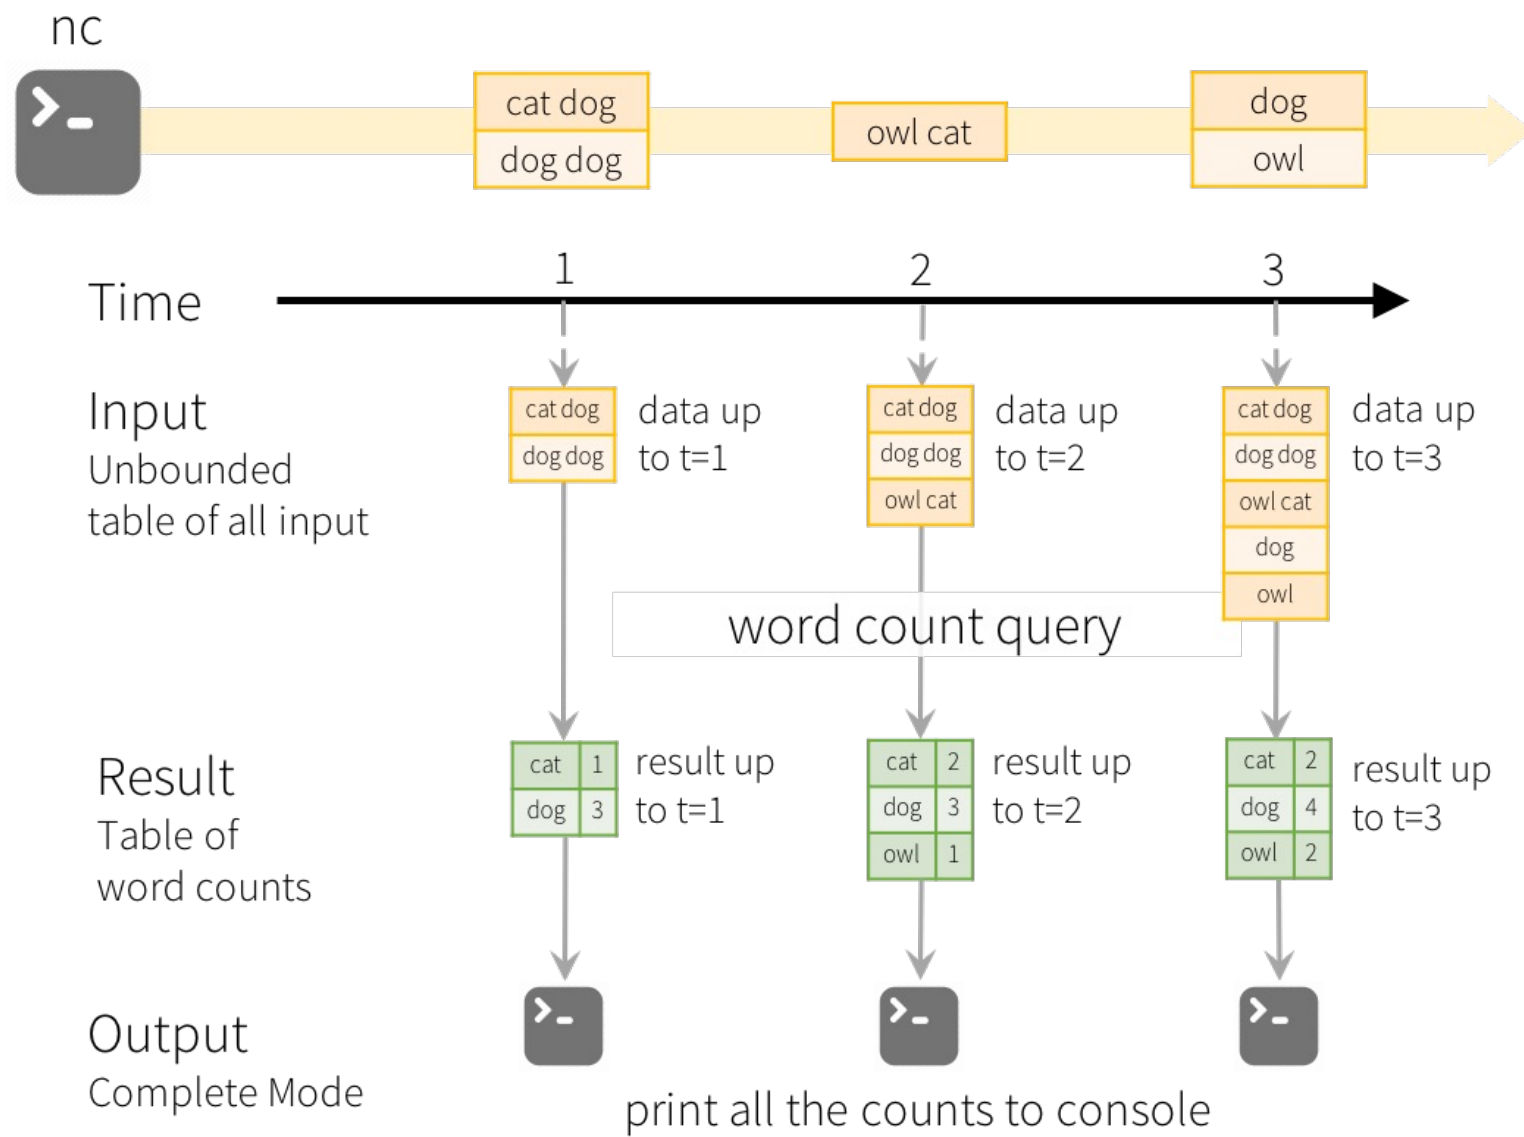
Programming Model for Structured Streaming

How does it work?


- Without changing the Dataset/DataFrame operations in queries,
- you will be able to choose the mode
- based on your application requirements.



Programming Model for Structured Streaming



Model of the Quick Example



Result Tables
after 5 minute triggers

12:00 - 12:10	cat	1
12:00 - 12:10	dog	3

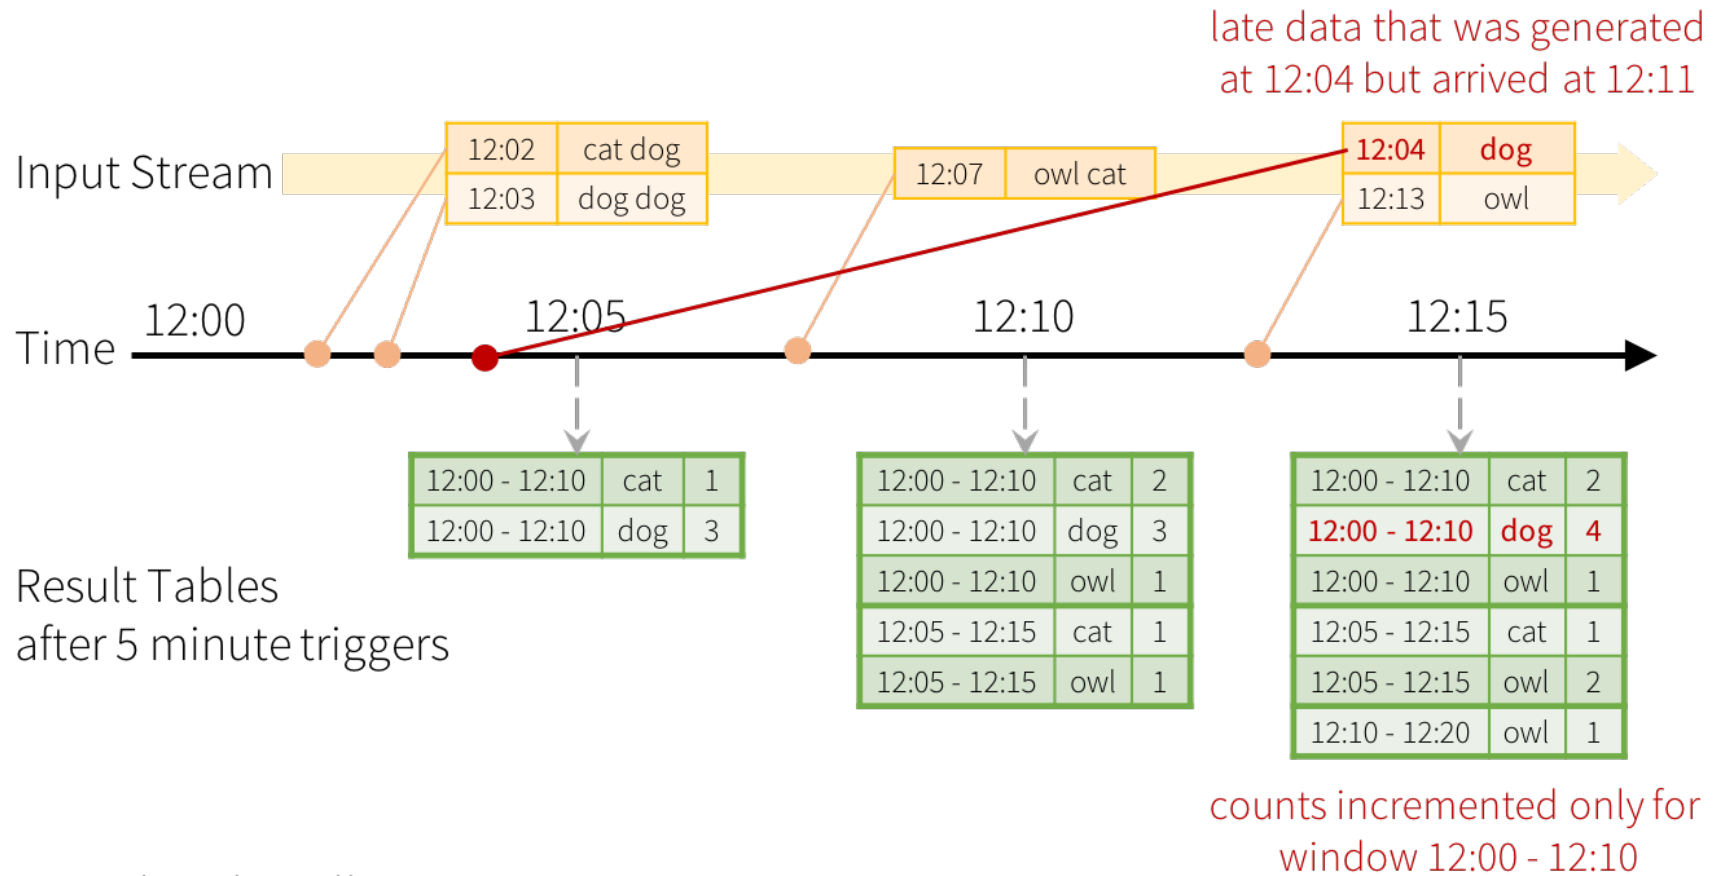
12:00 - 12:10	cat	2
12:00 - 12:10	dog	3
12:00 - 12:10	owl	1
12:05 - 12:15	cat	1
12:05 - 12:15	owl	1

counts incremented for windows
12:00 - 12:10 and 12:05 - 12:15

12:00 - 12:10	cat	2
12:00 - 12:10	dog	3
12:00 - 12:10	owl	1
12:05 - 12:15	cat	1
12:05 - 12:15	owl	2
12:05 - 12:15	dog	1
12:10 - 12:20	dog	1
12:10 - 12:20	owl	1

counts incremented for windows
12:05 - 12:15 and 12:10 - 12:20

Windowed Grouped Aggregation
with 10 min windows, sliding every 5 mins



Late data handling in
Windowed Grouped Aggregation

Steps for structured streaming



Let's say you want



to maintain a running word count



of text data received



from a data server



listening on a TCP socket.



Steps for structured streaming

- `import org.apache.spark.sql.functions._`
- `import org.apache.spark.sql.Session`



Steps for structured streaming

```
scala> val spark = SparkSession  
      .builder  
      .appName("StructuredNetworkWordCount")  
      .getOrCreate()
```

```
scala> import spark.implicits._
```



Steps for structured streaming

- `// Create DataFrame representing the stream of input lines from connection to localhost:9999`
- `val lines = spark.readStream`
- `.format("socket")`
- `.option("host", "localhost")`
- `.option("port", 9999)`
- `.load()`



Steps for structured streaming

- `// Split the lines into words`
- `val words = lines.as[String].flatMap(_.split(" "))`
- `// Generate running word count`
- `val wordCounts = words.groupBy("value").count()`



Steps for structured streaming

- `// Start running the query that prints the running counts to the console`
- `val query = wordCounts.writeStream`
- `.outputMode("complete")`
- `.format("console")`
- `.start()`

- `query.awaitTermination()`



Steps for structured streaming

- `$ nc -lk 9999`
 - apache spark
 - apache hadoop
-
- apache spark
 - apache hadoop

Steps for structured streaming

```
$ ./bin/run-example
```

```
org.apache.spark.examples.sql.streaming.StructuredNetworkWordCou  
nt localhost 9999
```

```
-----
```

```
Batch: 0
```

```
-----
```

```
+-----+-----+
```

```
| value|count|
```

```
+-----+-----+
```

```
| apache| 1|
```

```
| spark| 1|
```

```
+-----+-----+
```

Supported sources and sinks

- <https://spark.apache.org/docs/latest/structured-streaming-programming-guide.html>