

A silhouette of an industrial facility, likely a power plant or refinery, is visible against a warm, orange-hued sky at sunset or sunrise. Thick plumes of smoke or steam are rising from the facility, partially obscuring the sky. The foreground is dark and indistinct, suggesting a body of water or a flat landscape.

# Apache Spark SQL Optimization

Surendra Panpaliya



# Agenda

---

- Catalyst Optimizer
- What optimizations does catalyst do?
- "Explaining" Catalyst optimizations
- Overview of Explain output
- Optimization: Predicate Pushdown

# Agenda

---

Tungsten Optimizer

Tungsten Overview

Tungsten's Binary Format

# What is Catalyst?



Spark SQL was designed




with an optimizer



called Catalyst based on



the functional programming of Scala.



# What is Catalyst?



Main purposes are:



to add new optimization techniques



to solve some problems



with “big data”

# What is Catalyst?



Main purposes are:



2. to allow developers



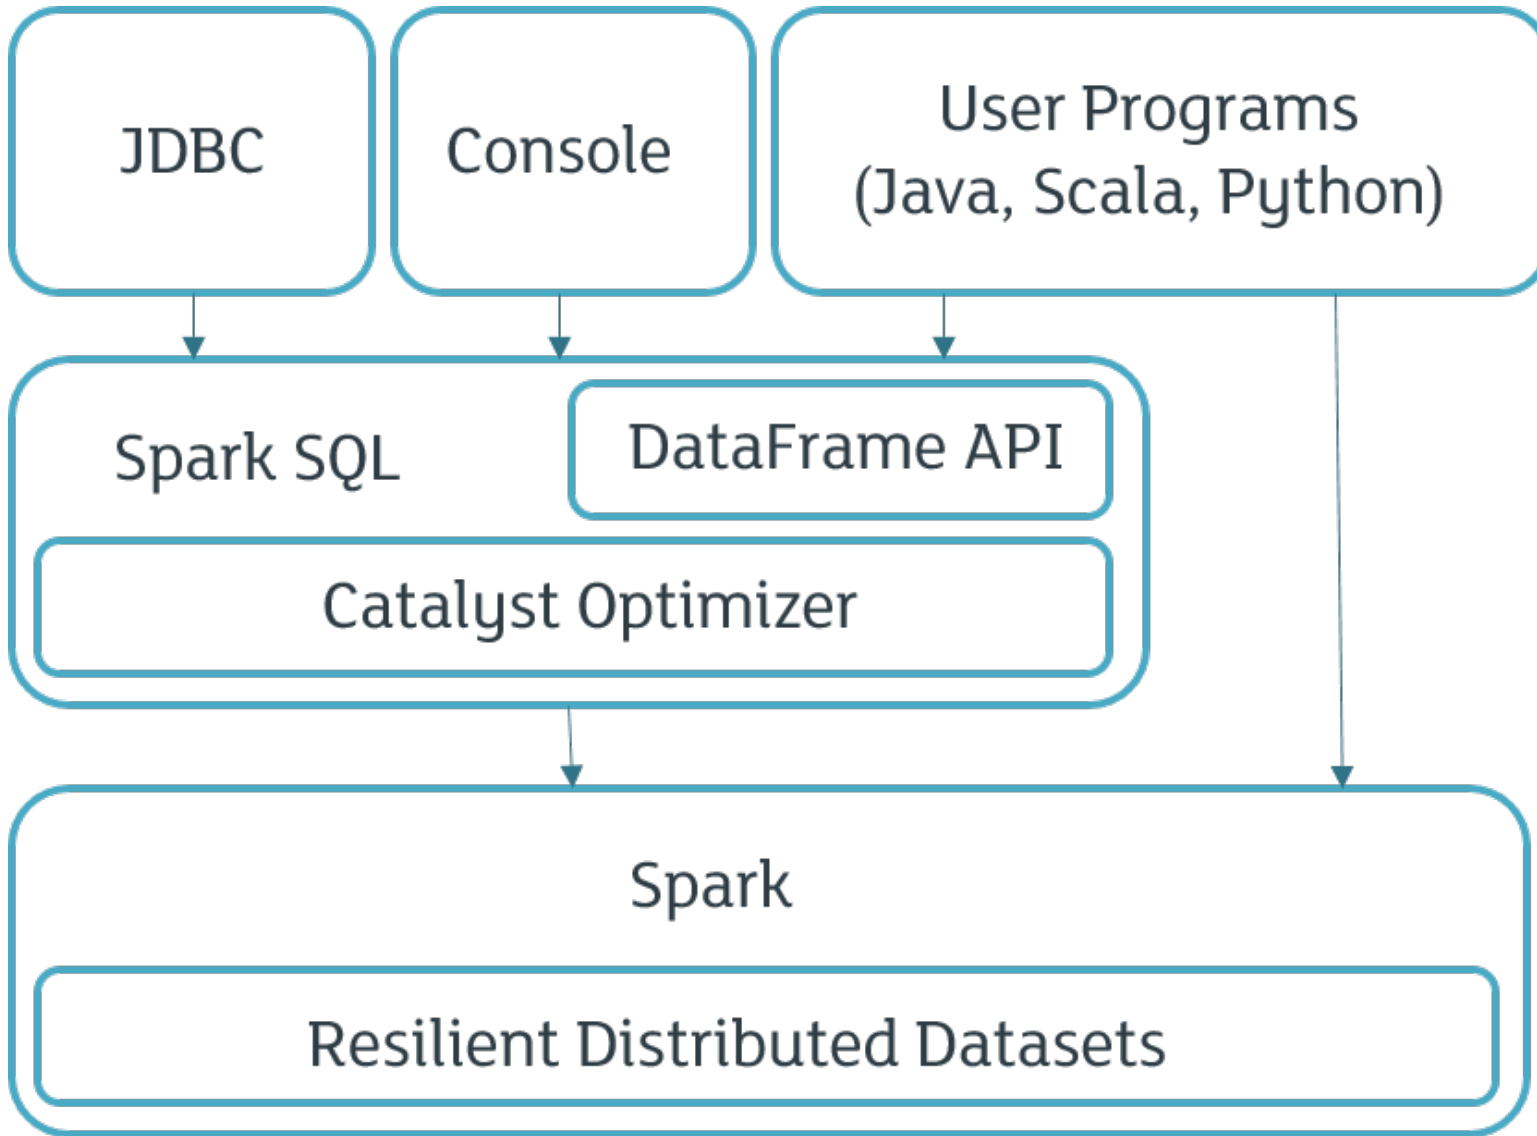
to expand and



customize the functions



of the optimizer.



Catalyst Spark  
SQL  
architecture  
and Catalyst  
optimizer  
integration



# Catalyst components

## Trees

The main data type in Catalyst is the tree.

Each tree is composed of nodes, and

each node has a node type and

zero or more children.





# Catalyst components

Trees

Objects are immutable and

can be manipulated

with functional language.

# Catalyst components

Trees

As an example

```
Merge(Attribute(x), Merge(Literal(1), Literal(2)))
```

# Catalyst components

Literal(value: Int): a constant value

Attribute(name: String): an attribute as input row

Merge(left: TreeNode, right: TreeNode):

mix of two expressions

# Rules

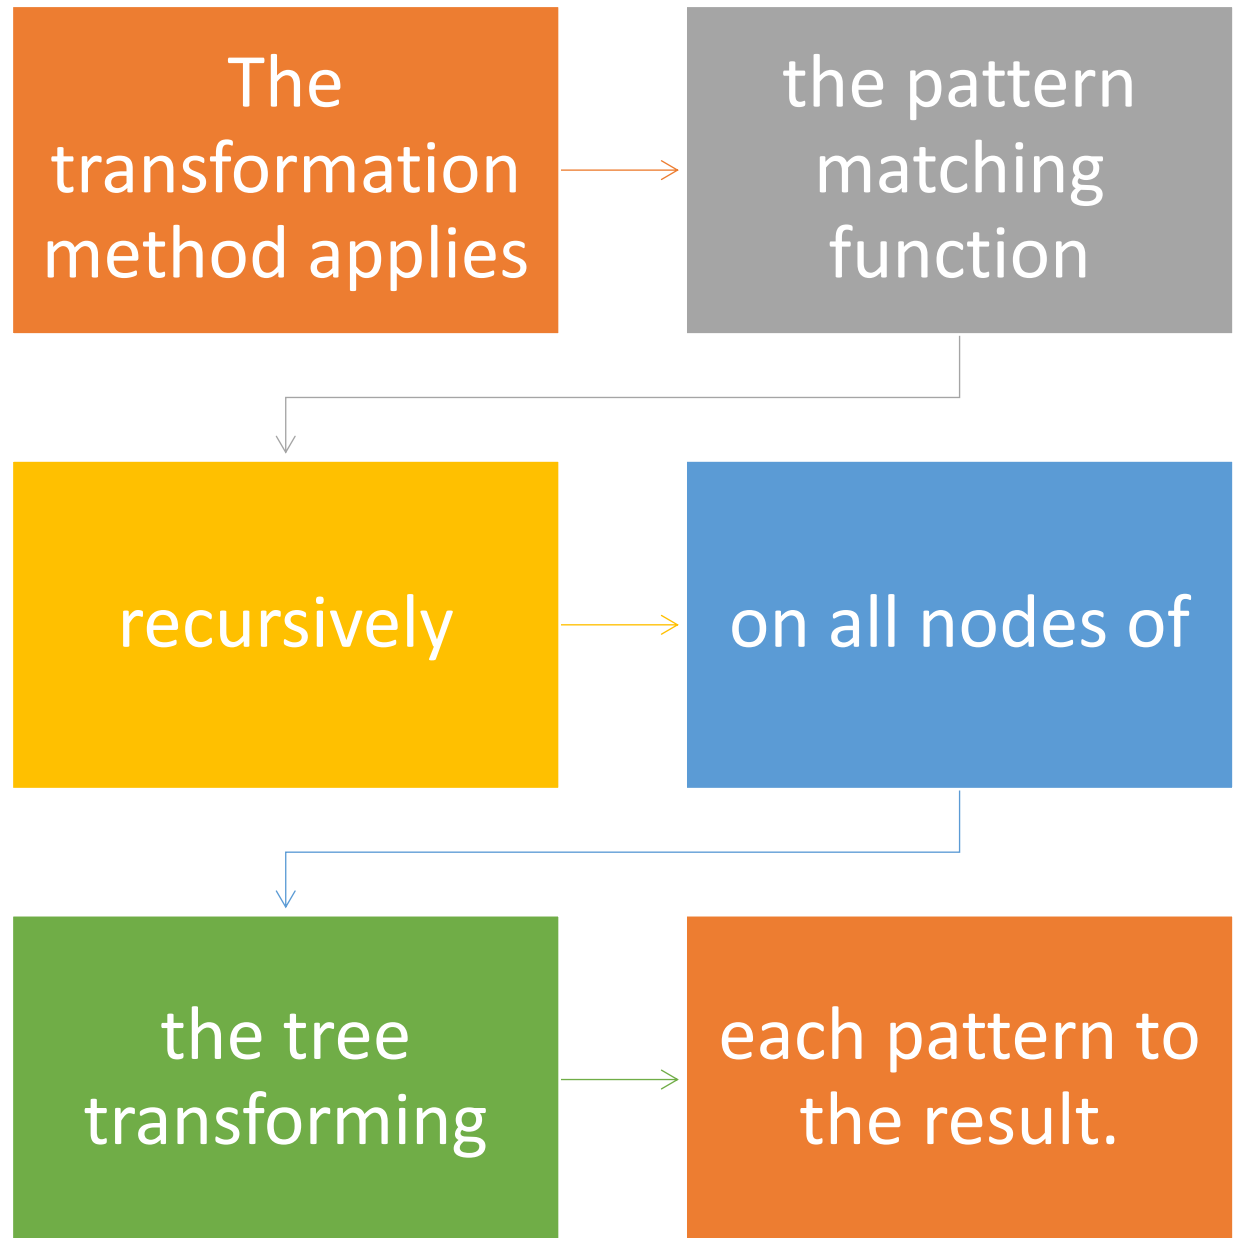
Trees can be manipulated

using rules,

which are functions of

a tree to another tree.

# Rules



# Rules

Example of a rule applied to a tree.

```
tree.transform {
```

```
  case Merge(Literal(c1), Literal(c2))  
    => Literal(c1) + Literal(c2)
```

```
}
```

# Using Catalyst in Spark SQL

The Catalyst Optimizer in Spark offers

rule-based and cost-based optimization.

Rule-based optimization indicates

how to execute the query

from a set of defined rules.

# Using Catalyst in Spark SQL

Cost-based  
optimization

generates multiple  
execution plans  
and

compares them to  
choose the lowest  
cost one.



# The Catalyst optimizer Example

- `case class EmployeeClass(empid: String, empname: String, salary: Int)`
- `val employeeDataset = Seq(
  - EmployeeClass("001", "Surendra", 4000000),
  - EmployeeClass("002", "Satish", 5000000)).toDS`
- `employeeDataset.show()`

# The Catalyst optimizer Example

- `val equery = employeeDataset.groupBy("salary").count().as("total")`
- `equery.show()`
- `equery.explain(extended = true)`

# Phases



**1. Analysis**



**2. Logic Optimization Plan**



**3. Physical plan**



**4. Code generation**

# 1. Analysis

The first phase of Spark SQL

optimization is the analysis.

Spark SQL starts

with a relationship

to be processed

that can be in two ways.

# 1. Analysis

A serious form from an

AST (abstract syntax tree)

returned by an SQL parser

DataFrame object

of the Spark SQL API.

## 2. Logic Optimization Plan



The second phase is the logical optimization plan.



In this phase, rule-based optimization is applied



to the logical plan.



It is possible to easily add new rules.

### 3. Physical plan

Spark SQL takes the logical plan and

generates one or more physical plans

using the physical operators

that match the Spark execution engine.

### 3. Physical plan

The plan to be executed is

selected using the cost-based model

(comparison between model costs).



# 4. Code generation



Code generation is the final phase of



optimizing Spark SQL.



To run on each machine,

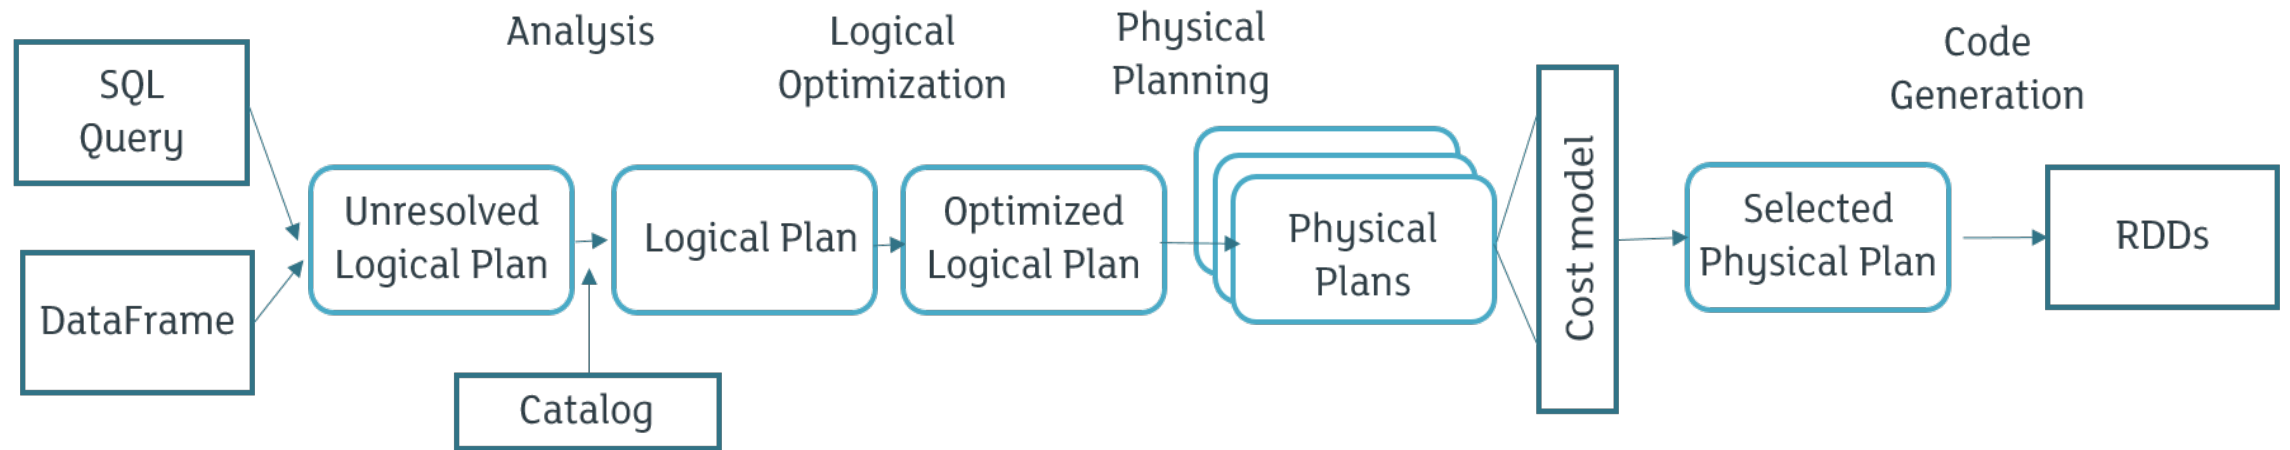


it is necessary to generate



Java code bytecode.

# *Phases of the query plan in Spark SQL. Rounded squares represent the Catalyst trees*



# What is Optimization in Apache Spark?

- Optimization is
- the method by which the long running
- application or system is fine tuned and
- made some change to make application
- to manage resources effectively and
- reduce the processing time efficiently.

# What is Optimization in Apache Spark?

- Apache Spark 2.0 is a major release
- that brought drastic changes
- to many framework,
- API's and libraries in that framework.
- In earlier version of Apache Spark,
- for optimization, developers need
- to change the source code of spark framework.

# What is Optimization in Apache Spark?

- This solved the problem
- for the particular user or
- to the particular account and
- also it is not advisable
- to change the code base of framework.

# What is Optimization in Apache Spark?

- There comes an idea
- to enhance the optimization features
- in later version that will act as
- an catalyst
- to run the application efficiently.

# What is Optimization in Apache Spark?

- In advanced Apache spark framework,
- we have a pluggable method
- which helps one
- to define a set of optimization rules and
- add it to the Catalyst.

# What is Optimization in Apache Spark?

- “The term **optimization** refers
- to a process in which
- a system is modified
- in such a way that
- it work more efficiently or
- it uses fewer resources.”



# What is Optimization in Apache Spark?

- [Spark SQL](#) is the most
- technically involved component of Apache Spark.
- Spark SQL deals with both
- SQL queries and DataFrame API.
- In the depth of Spark SQL
- there lies a **catalyst optimizer**.

# What is Optimization in Apache Spark?

- Catalyst optimization allows
- some advanced programming language
- features that allow
- you to build an extensible query optimizer.

# What is Optimization in Apache Spark?

- A new extensible optimizer called Catalyst
- emerged to implement Spark SQL.
- This optimizer is based on
- functional programming
- construct in Scala.

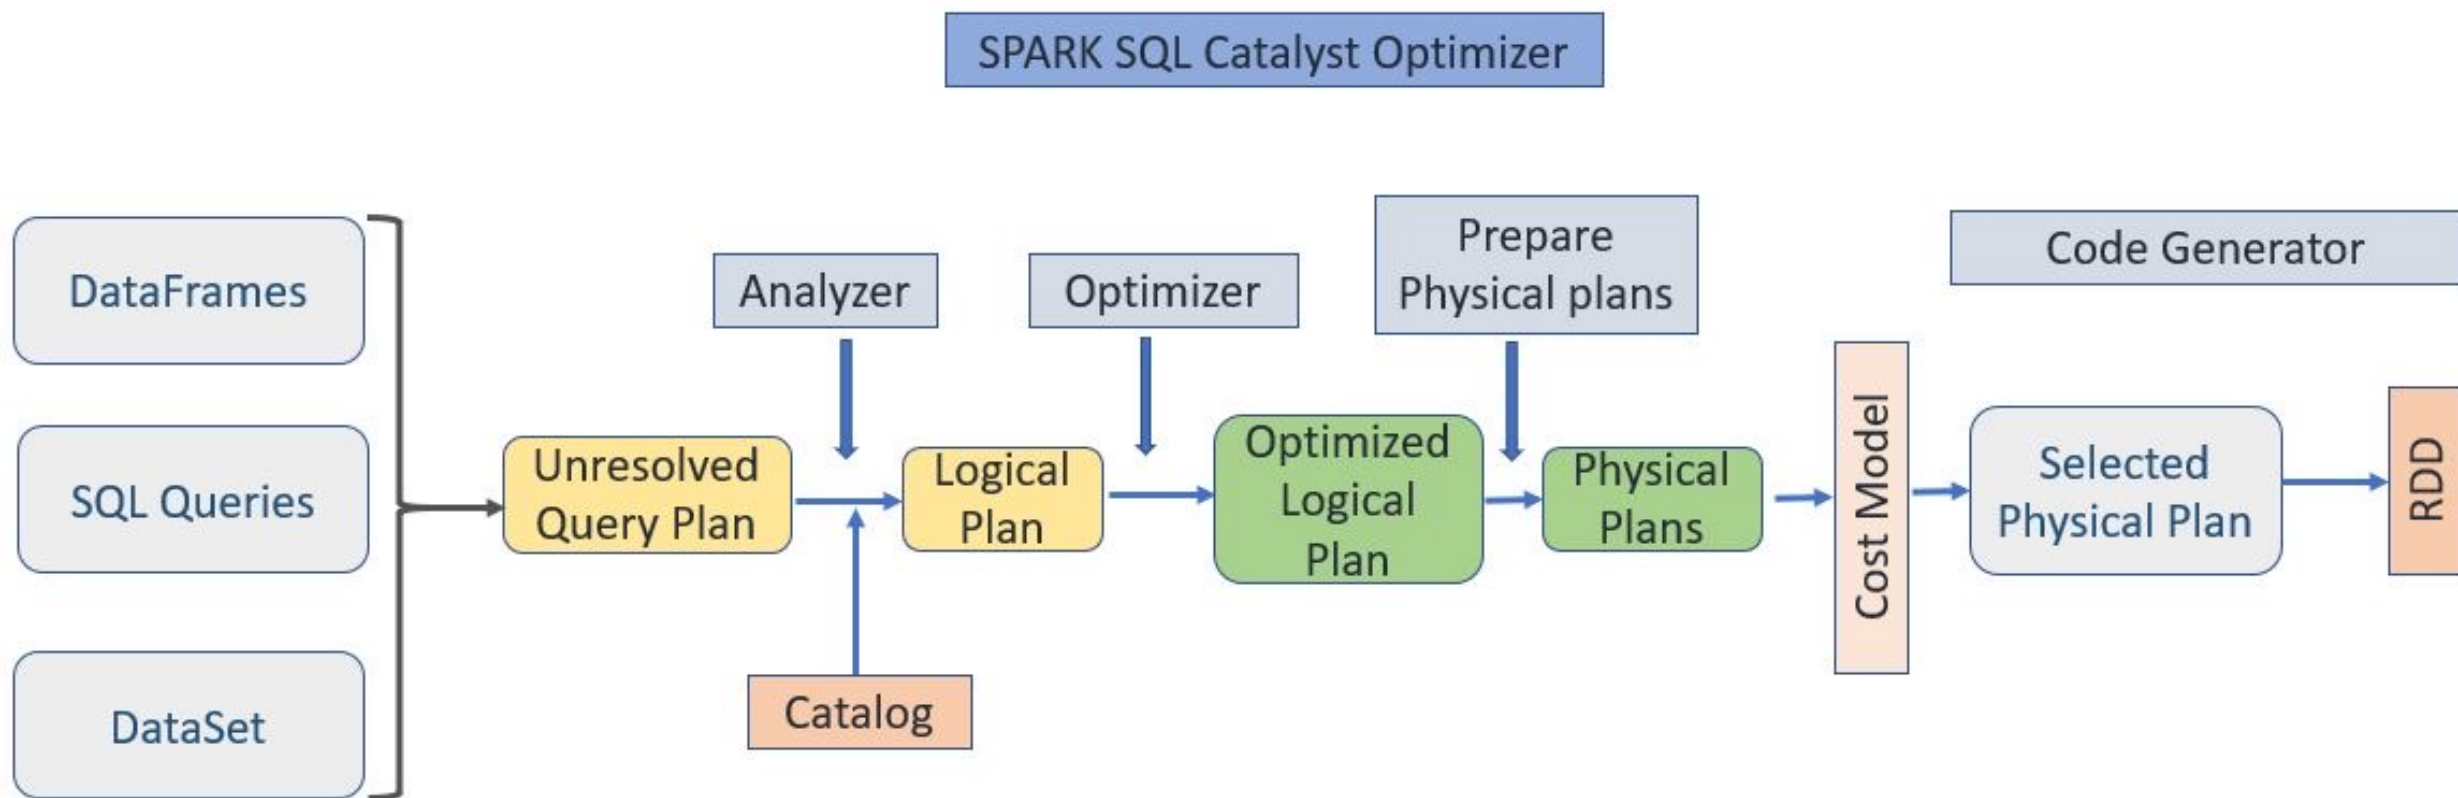
# What is Optimization in Apache Spark?

- Catalyst Optimizer supports both
- **rule-based** and **cost-based** optimization.
- In *rule-based optimization*
- the rule based optimizer
- use set of rule to determine
- how to execute the query.

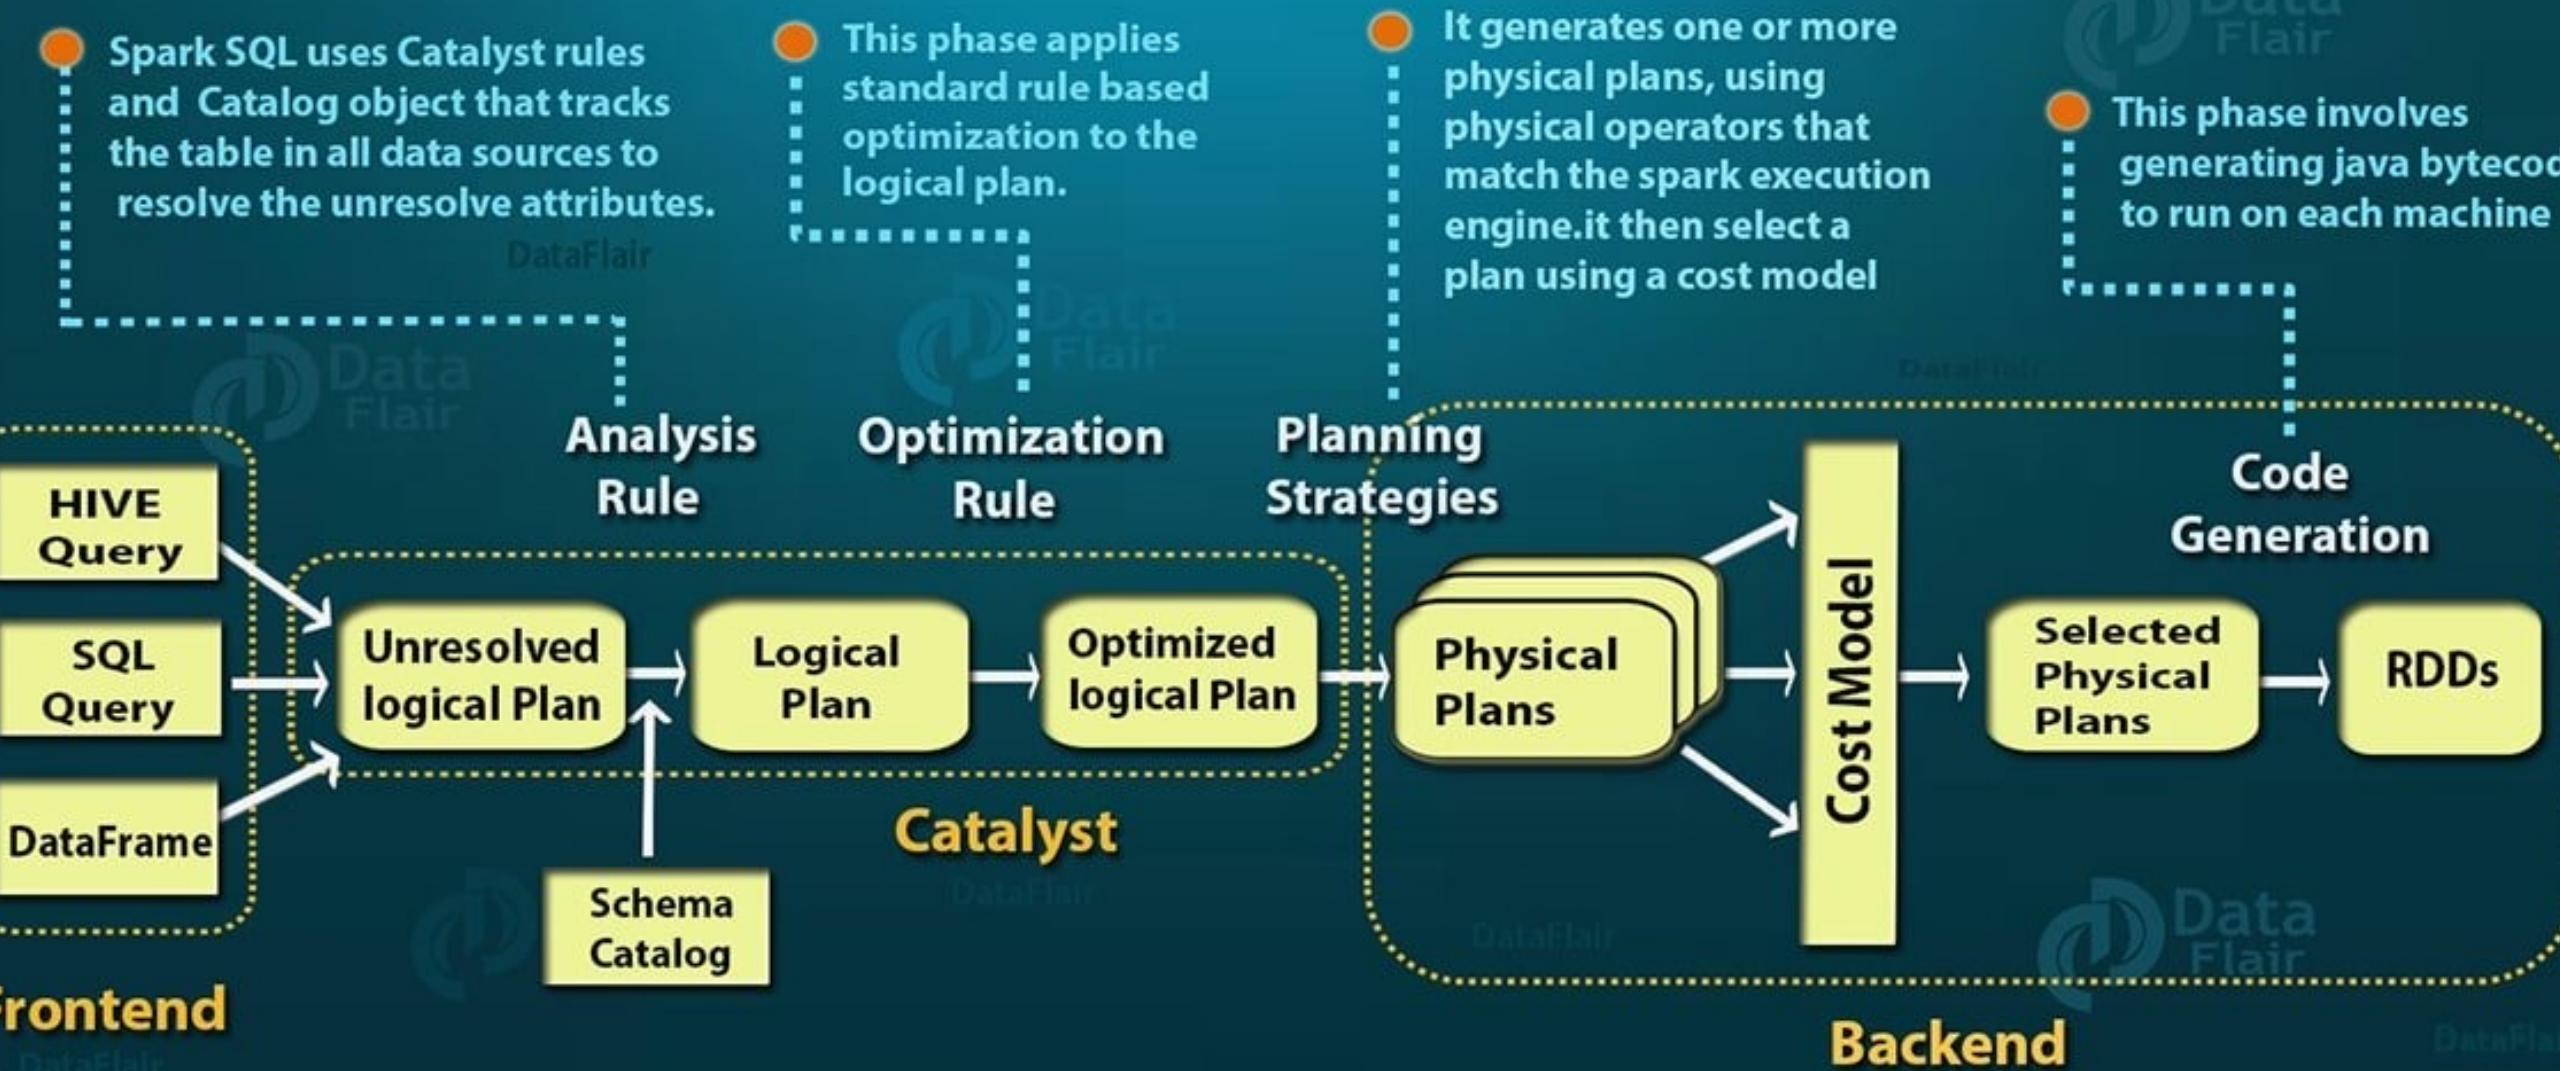
# What is Optimization in Apache Spark?

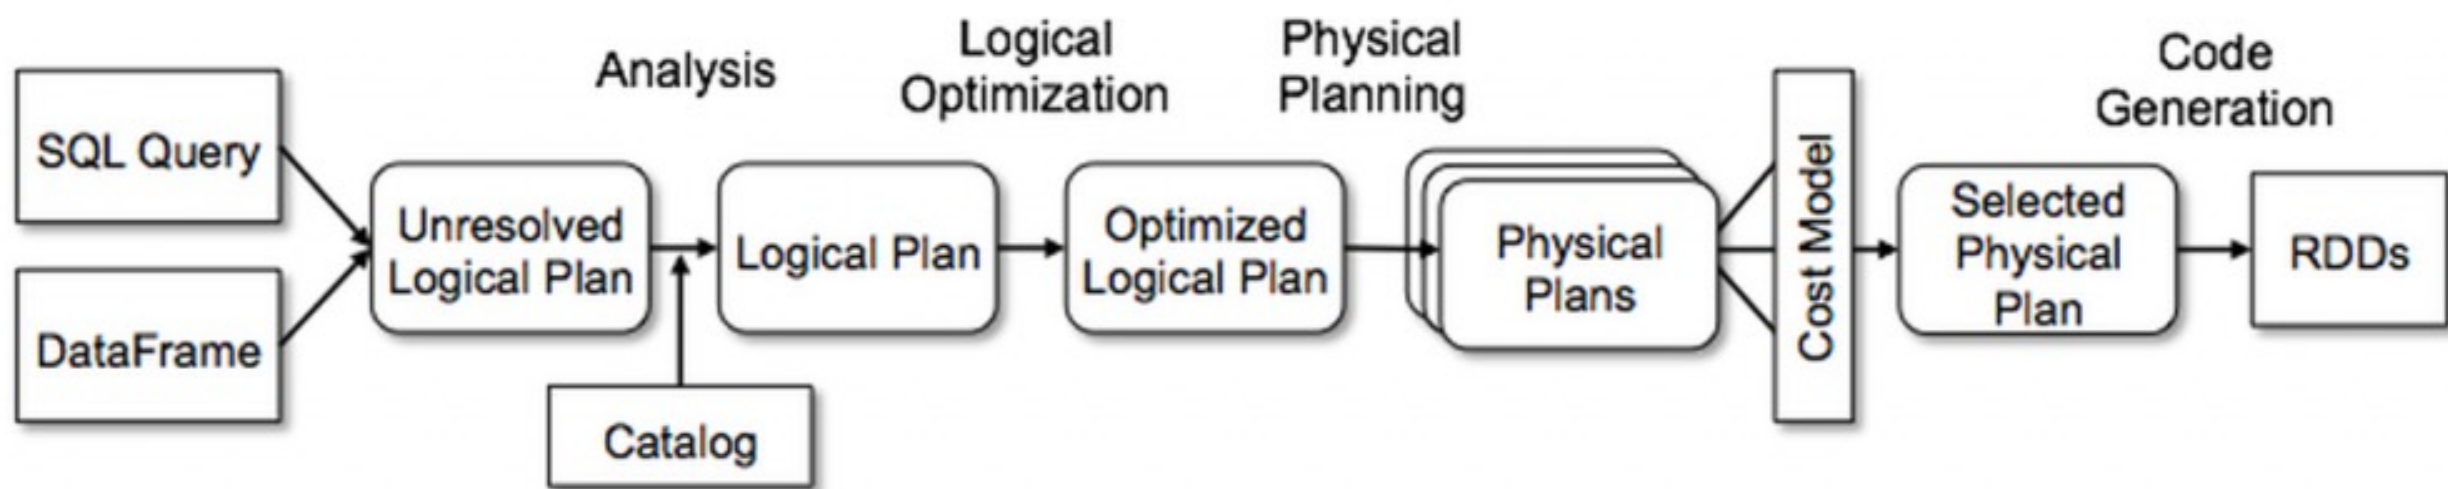
- While the *cost based optimization*
- finds the most suitable way
- to carry out SQL statement.
- In cost-based optimization,
- multiple plans are generated using
- rules and then their cost is computed.

# Execution Model



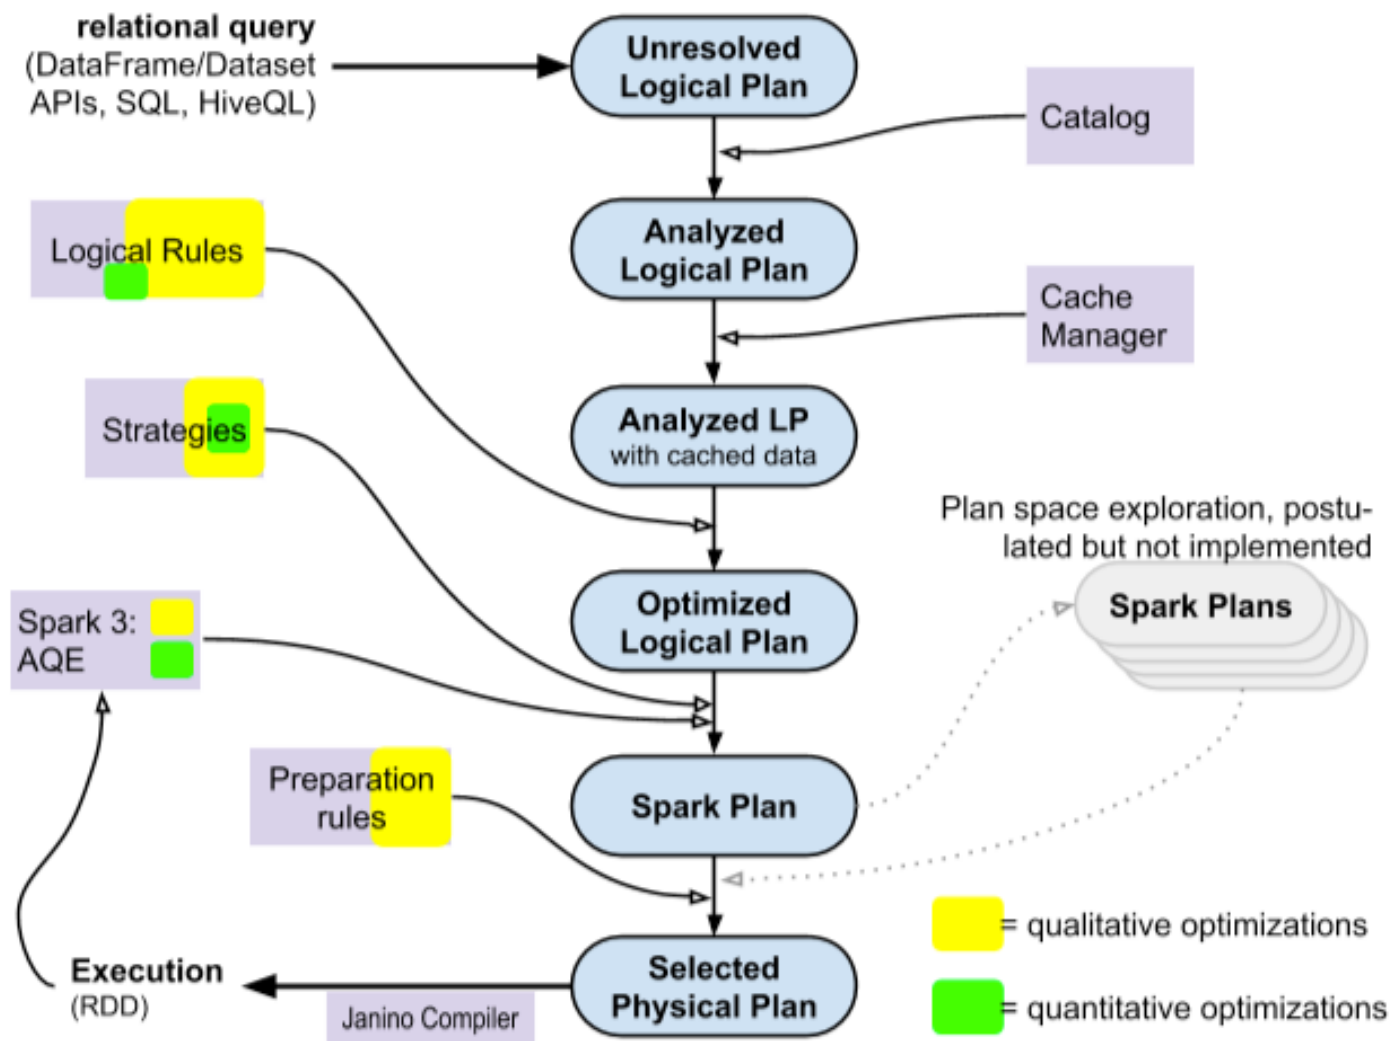
# Spark SQL Optimization







## The Catalyst Pipeline



# What is the need of Catalyst Optimizer?

- There are two purposes behind Catalyst's extensible design:
- We want to add the easy solution to tackle various problems with [Bigdata](#) like a problem with semi-structured data and advanced [data analytics](#).
- We want an easy way such that external developers can extend the optimizer.

# Fundamentals of Catalyst Optimizer

- Catalyst optimizer
- makes use of standard
- features of Scala programming like
- pattern matching.
- Catalyst contains the **tree** and
- the set of **rules**
- to manipulate the tree.

# Fundamentals of Catalyst Optimizer

- There are specific libraries
- to process relational queries.
- There are various rule sets
- which handle different
- phases of query execution

# Fundamentals of Catalyst Optimizer

- Like *analysis, query optimization,*
- *physical planning,* and
- *code generation*
- to compile parts of queries
- to Java bytecode.

## 4.1. Trees

- A tree is the main data type in the catalyst.
- A tree contains node object.
- For each node, there is a node.
- A node can have one or more children.

## 4.1. Trees

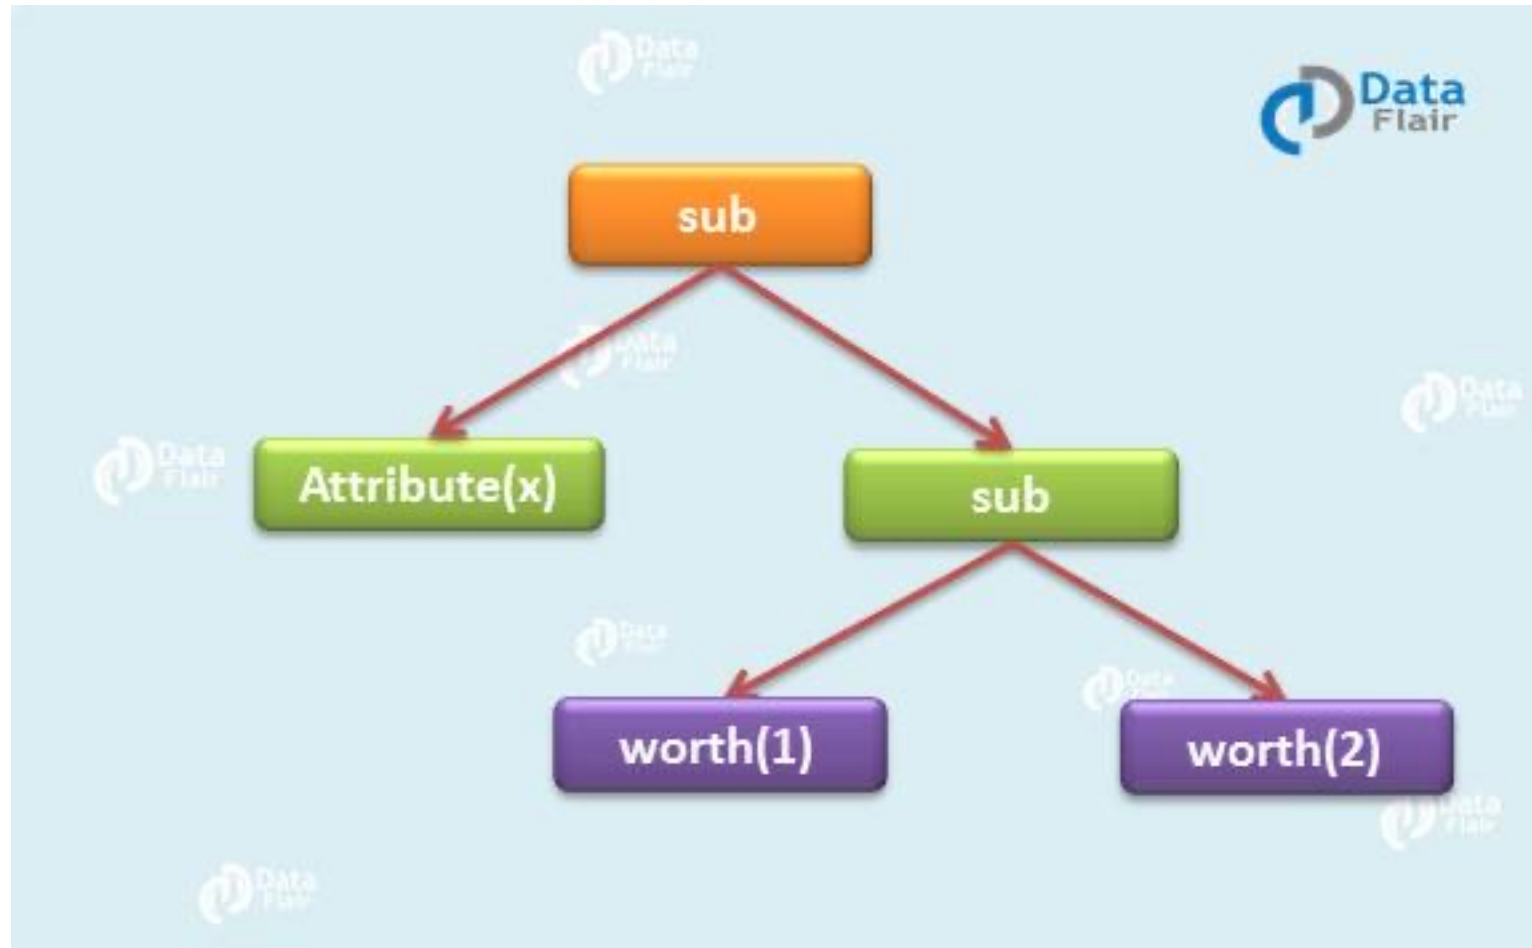
- New nodes are defined as
- subclasses of `TreeNode` class.
- These objects are immutable in nature.
- The objects can be manipulated
- using functional transformation.

# Trees

- For example, if we have three node classes:
- **worth**, **attribute**, and **sub** in which-
- `worth(value: Int)`: a constant value
- `attribute(name: String)`
- `sub (left: TreeNode, right: TreeNode)`:
- subtraction of two expressions.



# Trees



## 4.2. Rules

- We can manipulate tree using **rules**.
- We can define rules as a function from one tree to another tree.
- With rule we can run arbitrary code on input tree,
- the common approach
- to use a pattern matching function and
- replace subtree with a specific structure.

## 4.2. Rules

- In a tree with the help of
- **transform function**,
- we can recursively apply
- pattern matching on
- all the node of a tree.

## 4.2. Rules

- We get the pattern that
- matches each pattern
- to a result.

## 4.2. Rules

- For example:

*tree.transform {case Sub(worth(c1),worth(c2)) => worth(c1+c2) }*

## 4.2. Rules

- The expression that
- is passed during
- pattern matching
- to transform is a partial function.

## 4.2. Rules

- By partial function,
- it means it only needs to match
- to a subset of all possible input trees.

## 4.2. Rules

- Catalyst will see, to which part of a tree
- the given rule applies, and
- will automatically skip over
- the tree that does not match.
- With the same transform call,
- the rule can match multiple patterns.



## 4.2. Rules

- For example-
- *tree.transform {  
  case Sub(worth(c1), worth(c2)) => worth(c1-c2)  
  case Sub(left, worth(0)) => left  
  case Sub(worth(0), right) => right  
}*

## 4.2. Rules

- To fully transform a tree,
- rule may be needed
- to execute multiple time.

## 4.2. Rules

- Catalyst work by grouping rules
- into batches and
- these batches are executed
- until a fixed point is achieved.

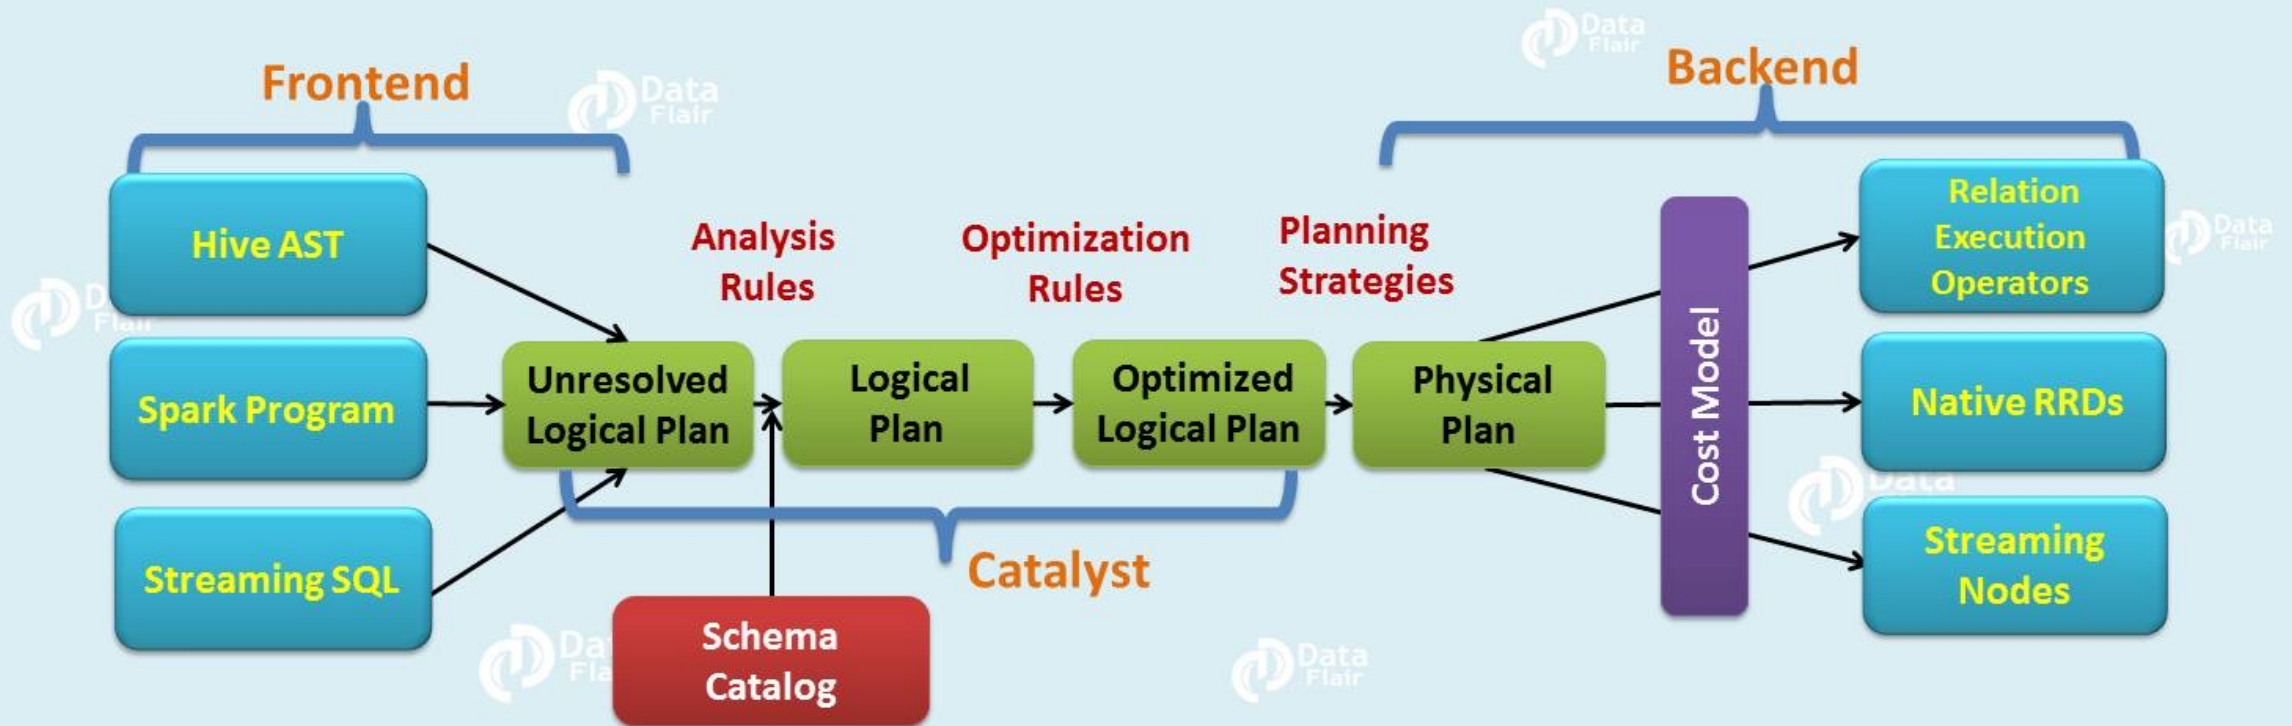
## 4.2. Rules

- Fixed point is a point
- after which
- tree stops changing
- even after applying rules.

## 5. Spark SQL Execution Plan

- In four phases we use Catalyst's general tree transformation framework:
- Analysis
- Logical Optimization
- Physical planning
- Code generation

# Spark SQL Execution Plan



## 5.1. Analysis

- Spark SQL Optimization starts from relation to be computed. It is computed either from **abstract syntax tree (AST)** returned by **SQL parser** or [dataframe](#) object created using API.
- Both may contain unresolved attribute references or relations.

## 5.1. Analysis

- By unresolved attribute,
- it means we don't know its type or
- have not matched it to an input table.



## 5.1. Analysis

- Spark SQL make use of Catalyst rules and
- a Catalog object that track data
- in all data sources
- to resolve these attributes.

## 5.1. Analysis

- It starts by creating
- an unresolved logical plan, and
- then apply the following steps:
- Search relation BY NAME FROM CATALOG.
- Map the name attribute

## 5.1. Analysis

- For example,
- col, to the input provided
- given operator's children.

## 5.1. Analysis

- Determine which attributes
- match to the same value
- to give them unique ID.
- Propagate and push type
- through expressions

# Logical Optimization

- In this phase of Spark SQL optimization,
- the standard rule-based optimization is applied to the logical plan.
- It includes **constant folding**, **predicate pushdown**, **projection pruning** and other rules.
- It became extremely easy
- to add a rule for various situations.

# Physical Planning

- There are about 500 lines of code in the physical planning rules.
- In this phase, one or more physical plan is formed
- from the logical plan,
- using physical operator matches
- the Spark execution engine.
- And it selects the plan using the cost model.

# Physical Planning

- It uses Cost-based optimization only to select join algorithms.
- For small relation SQL uses broadcast join,
- the framework supports broader
- use of cost-based optimization.
- It can estimate the cost recursively
- for the whole tree using the rule.

# Physical Planning

- Rule-based physical optimization,
- such as pipelining projections or filters
- into one Spark *map* Operation is also carried out
- by the physical planner.
- Apart from this, it can also push operations from the logical plan into data sources that support predicate or projection pushdown.



# Code Generation

- The final phase of Spark SQL optimization is code generation.
- It involves generating Java bytecode to run on each machine.
- Catalyst uses the special feature of Scala language,
- “**Quasiquotes**” to make code generation easier
- because it is very tough to build code generation engines.

# Code Generation

- Quasiquotes lets the programmatic construction of abstract syntax trees (ASTs) in the Scala language,
- which can then be fed to the Scala compiler at runtime to generate bytecode.
- With the help of a catalyst,

# Code Generation

- we can transform
- a tree representing an expression in SQL
- to an AST for Scala code
- to evaluate that expression, and
- then compile and
- run the generated code.

# Conclusion

- Spark SQL optimization
- enhances the productivity of developers and
- the performance of the queries that they write.

# Conclusion

- A good query optimizer
- automatically rewrites relational queries
- to execute more efficiently

# Conclusion

- Uses techniques such as
- filtering data early,
- utilizing available indexes, and
- even ensuring different data sources are
- joined in the most efficient order

# Conclusion

- By performing these transformations,
- the optimizer improves the execution times
- of relational queries and
- frees the developer from
- focusing on the semantics of their application
- instead of its performance.